

3.2 Modeling

Heuristics, classical models, and when to stack them.

1. Start with a rule you can write today

Modeling does not begin with a neural net. It begins with a decision you can ship.

Heuristics are rules a person wrote. A spam filter may already have a blacklist of domains and a list of words that almost always mean junk. Those rules are a model. They have precision you can measure. They also fail in the usual way: language moves, and a pile of rules becomes a system nobody can debug.

Use heuristics when you have little labeled data, when a stakeholder needs a first version, or when a pattern is cheap and stable (`unsubscribe` in the subject line).

Existing APIs (a vendor toxicity score, a cloud language API) are the same idea: a model you did not train, used as a baseline.

2. Then fit a model, and keep the rules as features

As you collect labels, a learned model usually beats a growing rule pile. Do not throw the rules away. Two honest ways to keep them:

Make a feature. The blacklist hit becomes a 0/1 column next to TF-IDF. The model learns how much to trust it.

Pre-process the input. Normalize URLs, mask emails, drop the signature block, then send the cleaned text to the classifier. The heuristic is now part of the pipeline, not a competing system.


If you keep adding `if` statements after the model, you are back to an unmaintainable rule base. Put the knowledge in features or in the input, then retrain.

3. Building “the” model is iteration

A first baseline (heuristic or Naive Bayes on counts) is required. Then you iterate until the thing is accurate *and* operable.

Lever	What you change
Better features	Note 3.1 : new counts, n-grams, embeddings
Ensemble / stacking	Combine several models
Transfer learning	Start from weights trained on a larger corpus
Re-apply heuristics	Catch the remaining systematic errors

Stacking: train models that disagree in useful ways (a linear bag-of-words model and a small neural net). A second-stage model takes their scores as input and makes the call.



Transfer learning is the usual path once you leave classical features: a pre-trained language model, fine-tuned on your labels. It is still a modeling choice, not magic. If your data is tiny and unlike the pre-training domain, a linear model on good features can win.

4. Let the data pick the path

Volume and quality decide the strategy more than taste.

Data you have	Sensible first model
Almost no labels, clear rules	Heuristics or a vendor API
Hundreds of clean labels	Linear model on handcrafted or TF-IDF features
Thousands of in-domain labels	Classical ML or a small neural net
Large labeled set, messy text	Deep model / transfer learning
Lots of unlabeled text, few labels	Pre-train or use the unlabeled text for embeddings; label what the model is unsure of

Data attribute	Decision path	Examples
Large data volume	Can use techniques that require more data, like DL. Can use a richer set of features as well. If the data is sufficiently large but unlabeled, we can also apply unsupervised techniques.	If we have a lot of reviews and metadata associated with them, we can build a sentiment-analysis tool from scratch.
Small data volume	Need to start with rule-based or traditional ML solutions that are less data hungry. Can also adapt cloud APIs and generate more data with weak supervision. We can also use transfer learning if there's a similar task that has large data.	This often happens at the start of a completely new project.
Data quality is poor and the data is heterogeneous in nature	More data cleaning and pre-processing might be required.	This entails issues like code mixing (different languages being mixed in the same sentence), unconventional language, transliteration, or noise (like social media text).
Data quality is good	Can directly apply off-the-shelf algorithms or cloud APIs more easily.	Legal text or newspapers.
Data consists of full-length documents	Choose the right strategy for breaking the document into lower levels, like paragraphs, sentences, or phrases, depending on the problem.	Document classification, review analysis, etc.

"Production-ready" is part of the choice. A 2-point accuracy gain that cannot meet the latency budget is not an improvement. Note **3.4** is that constraint.

5. What "baseline" means here

Write the dummy first: majority class, or "always sales," or "keyword `refund` → care." Report its F1. Every later model has to beat *that* number, not a number from a blog.

Then one classical model (Naive Bayes or logistic regression on TF-IDF). Then, if you still need more and you have the data, an ensemble or a transferred net.

If the jump from dummy to TF-IDF is huge and the jump from TF-IDF to a transformer is 0.4 F1, stop. The labels or the extract are the problem (notes **2.1–2.2**), not the architecture.

Keep a sheet: data size, features, model, intrinsic scores, latency. That sheet is the model. The `.pk1` file is just the current row.

6. Practice

1. Write one heuristic for the sales-vs-care router from note **2.1**. Then say how you would turn it into a feature.
2. You have 400 labeled tickets and 200,000 unlabeled ones. What modeling path from the table do you take first?
3. Why can stacking two weak models beat either one, and when does it not?