

14.4 Text Summarization

Extractive versus abstractive shortening of a document.

1. Shorten the document, keep the point

Text summarization turns a long document into a shorter text that a person (or an indexer) can use. The goal is not "as short as possible." It is a coherent piece that still carries the main ideas.

Benefits you can sell: faster reading, a smaller search index, a cleaner store of "what this file was about."

This note follows the same applied texts as the rest of the course: *Practical Natural Language Processing* and *NLP in Action*.

2. Three design choices

Research since the DUC conferences organizes the job along three axes.

Axis	Option A	Option B
How the words are produced	Extractive: pick sentences that already exist	Abstractive: write new sentences
Whether a query exists	Query-focused: answer this information need	Query-independent: a general digest
How many sources	Single-document	Multi-document: fuse a cluster

Extractive is copy-and-rank. It cannot invent a fact, and it can sound choppy. Abstractive can read more naturally and can also hallucinate.


Query-focused summaries change with the user. Query-independent summaries are what you put at the top of an article or in an index.

Multi-document work has to handle repetition and contradiction across files. Do not start there.

3. What people actually ship

The common industrial case is **single-document, query-independent, extractive**.

Purpose: a short digest for a reader, or a stand-in document for a search index.



Two other patterns from real deployments:

- **Highlight instead of rewrite.** Color the sentences that carry the gist of a news story. The user still sees the article. You only need a sentence ranker.
- **Index the summary.** Store the digest, not the full text, when the index is too large. Retrieval quality then depends on the digest.

Both are extractive problems with a UI choice on top.

4. Practical advice

Pre-processing dominates. Sentence splitting and HTML cleanup decide what the ranker is allowed to pick. Library splitters fail on lists, captions, and "U.S." A bad split becomes a bad "sentence" in the summary. Write a splitter for your format if the default is clearly wrong.

Length hurts some algorithms. Graph methods (TextRank and cousins) get expensive on long inputs. Two workarounds:

- summarize chunks, then summarize the union
- take the first (M\%) and last (N\%) when those parts hold the thesis (same trick as KPE)

Evaluate with a task. ROUGE against a reference is the research default. In a product, measure whether people click, whether search still finds the file, or whether a highlighter agrees with an editor. A fluent abstractive paragraph that adds a date the article never stated is a failure.

Lead and conclusion bias is real: extractive rankers love the first paragraph. That is often correct for news and wrong for a methods section. If your corpus is scientific or legal, do not trust a "first 3 sentences" baseline without checking.

Graph summarizers share DNA with graph KPE (note **10.3**): sentences are nodes, overlap is an edge, PageRank picks the summary. The same length and overlap warnings apply.

5. What this note is for

You should be able to:

- place a project on the three axes above
 - say why extractive is the default first system
 - name the two production failure modes: bad sentence splits and too-long graphs
-

6. Practice

1. A legal tool must not invent a clause. Extractive or abstractive? Why?
2. You want "sentences that answer *What did the Fed do?*" Which cell of the table is that?
3. TextRank times out on 80-page filings. What do you change before you change the algorithm?

Chunk first. A summary of summaries is cheaper than a graph on 2,000 sentences.