

6.2 Traditional Text Classification

BoW or TF-IDF features with a linear or neighbor model.

1. Two classical classifiers


Once the document is a vector (BoW or TF-IDF from Week 4), you can use any supervised model. Lecture walks two that you should be able to explain: **(k)-nearest neighbors** and a **linear SVM**.

The rest of the note is how you **choose hyperparameters** and how you **score** the classifier when classes are not 50/50. That part matters more than the algebra.

Practical Natural Language Processing is the course text. The pictures of CV and metrics are standard ML.

2. (k)-nearest neighbors

Store the training set. To label a new point, find the (k) nearest training vectors (Euclidean or cosine) and take a **majority vote**.




There are no weights to learn. The whole training matrix is the model. That is fine for a few thousand short documents and painful for a million.

(k) is a hyperparameter. Small (k) follows local noise. Large (k) smooths toward the majority class. On TF-IDF, **cosine** is the usual distance.

3. Linear SVM

Training pairs $(\mathbf{x}_i, y_i)$ with $(y_i \in \{+1, -1\})$. A linear SVM finds a hyperplane $(\mathbf{w} \cdot \mathbf{x} - b = 0)$ that separates the classes with the **largest margin**.

You minimize $(\|\mathbf{w}\|)$ subject to $(y_i(\mathbf{w} \cdot \mathbf{x}_i - b) \geq 1)$.



In text, $(\mathbf{x})$ is usually L2-normalized TF-IDF. `LinearSVC` or `SGDClassifier(loss="hinge")` is what you call. A kernel SVM on a 50,000-column bag is rarely worth it. The linear version is the baseline to beat.

4. Cross-validation and hyperparameters

Parameters are what the learner sets from data: $(\mathbf{w})$, the support vectors, the neighbor labels.

Hyperparameters are what you set: (k) , the SVM's (C) , `max_features`, `n-gram range`. If you tune (C) on the same fold you report, the score is optimistic.

Hold-out: train / validation / test. Fit on train, pick hyperparameters on validation, report once on test.

(K)-fold: split train into (K) pieces. Each piece is validation once. Average the (K) scores. Use this when the data set is not huge.

Nested CV: the inner loop picks hyperparameters; the outer loop estimates generalization. That is the honest number when you need one number for a paper or a client and you already used the data to search.


If you only remember one rule: the test set is not a tuning knob.

5. Metrics when accuracy lies

Accuracy is correct / total. It is fine when classes are balanced. It is not fine when 95% of tickets are "general inquiry." A model that always says "general" is 95% accurate and useless.

For a class you care about (say *happy*):

Metric	Question it answers
Precision	Of the docs I called <i>happy</i> , how many were?
Recall	Of the <i>happy</i> docs, how many did I find?
F1	Harmonic mean of the two
ROC	True-positive rate vs false-positive rate as you move the threshold



F1 punishes a system that only wins on one of the two. ROC / AUC is useful when you will pick a threshold later.

sklearn: `precision_score`, `recall_score`, `f1_score`, `roc_curve`, and `classification_report`.

```
from sklearn.metrics import classification_report
print(classification_report(y_true, y_pred, target_names=["class 0", "class 1", "class 2"]))
```

Read the per-class rows, not just the headline accuracy.

6. Practice

1. Training data are 90% ham. A filter that never flags spam has high accuracy. Which metric exposes that it found no spam?

2. You pick (k) by maximizing accuracy on the test set. What is wrong, and what should you use instead?
3. Why is a **linear** SVM the usual text baseline rather than an RBF kernel SVM?