Introduction to ADMM and Penalty Decomposition for Intersection Constraints

Introduction to ADMM

The Alternating Direction Method of Multipliers (ADMM) is an optimization algorithm designed to solve problems with separable objectives and constraints. It is particularly effective for large-scale and distributed optimization problems and is widely used in fields like machine learning, signal processing, and statistics.

Generic Problem Formulation

ADMM solves problems of the form:

$$\min_{x,z} f(x) + g(z) \quad \text{s.t.} \quad Ax + Bz = c,$$

where:

- $f(x)$ and $g(z)$ are convex functions,
- $A, B,$ and c are problem-specific matrices/vectors.

Augmented Lagrangian

To handle the constraint $Ax + Bz = c$, ADMM introduces the augmented Lagrangian:

$$L_\rho(x, z, \lambda) = f(x) + g(z) + \lambda^T(Ax + Bz - c) + \frac{\rho}{2}\|Ax + Bz - c\|_2^2,$$

where:

- λ is the dual variable (Lagrange multiplier),
- $\rho > 0$ is the penalty parameter.

ADMM Iterative Updates

ADMM alternates between optimizing x, z, and updating λ. At iteration k:

1. **x-Update**:
$$x^{k+1} := \arg\min_x \left(f(x) + \frac{\rho}{2}\|Ax + Bz^k - c + \frac{\lambda^k}{\rho}\|_2^2 \right).$$
2. **z-Update**:
$$z^{k+1} := \arg\min_z \left(g(z) + \frac{\rho}{2}\|Ax^{k+1} + Bz - c + \frac{\lambda^k}{\rho}\|_2^2 \right).$$
3. **Dual Variable Update**:
$$\lambda^{k+1} := \lambda^k + \rho(Ax^{k+1} + Bz^{k+1} - c).$$

Reformulating Intersection Constraints in ADMM

Consider an optimization problem where the feasible set is an intersection of two sets:

$$\min_{x \in X \cap Z} f(x) + g(x).$$

The constraint $x \in X \cap Z$ can be reformulated by introducing an auxiliary variable z:

$$\min_{x,z} f(x) + g(z) \quad \text{s.t.} \quad x \in X,\, z \in Z,\, x - z = 0.$$

Augmented Lagrangian for Reformulation

The augmented Lagrangian becomes:

$$L_\rho(x, z, \lambda) = f(x) + g(z) + \lambda^T (x - z) + \frac{\rho}{2}\|x - z\|_2^2.$$

ADMM Iterative Updates

1. x**-Update**:
$$x^{k+1} := \arg\min_{x \in X} \left(f(x) + \frac{\rho}{2}\|x - z^k + \frac{\lambda^k}{\rho}\|_2^2 \right).$$
2. z**-Update**:
$$z^{k+1} := \arg\min_{z \in Z} \left(g(z) + \frac{\rho}{2}\|x^{k+1} - z + \frac{\lambda^k}{\rho}\|_2^2 \right).$$
3. **Dual Variable Update**:
$$\lambda^{k+1} := \lambda^k + \rho(x^{k+1} - z^{k+1}).$$

Connection to Penalty Decomposition (PD)

In penalty decomposition, the constraint $x - z = 0$ is enforced using a penalty term without introducing a dual variable:

$$\min_{x,z} f(x) + g(z) + \frac{\rho}{2}\|x - z\|_2^2.$$

PD Updates

1. x**-Update**:
$$x^{k+1} := \arg\min_{x \in X} \left(f(x) + \frac{\rho}{2}\|x - z^k\|_2^2 \right).$$
2. z**-Update**:
$$z^{k+1} := \arg\min_{z \in Z} \left(g(z) + \frac{\rho}{2}\|x^{k+1} - z\|_2^2 \right).$$

Improved Example: Sparse Recovery with Simpler Subproblems

Consider the sparse regression problem:

$$\min_{x \in \mathbb{R}^n} \|x\|_1 + \frac{\gamma}{2}\|Ax - b\|_2^2 \quad \text{s.t.} \quad x \geq 0.$$

which can be rewritten as follows:

$$\min_{x \in \mathbb{R}^n, z \in \mathbb{R}^n} \|x\|_1 + \frac{\gamma}{2}\|Ax - b\|_2^2 \quad \text{s.t.} \quad x - z = 0,\, z \geq 0.$$

ADMM Updates

1. **x-Update** (soft-thresholding for sparsity):

$$x^{k+1} := \arg\min_x \left(\|x\|_1 + \frac{\rho}{2}\|x - z^k + \frac{\lambda^k}{\rho}\|_2^2 \right).$$

2. **z-Update** (non-negativity constraint):

$$z^{k+1} := \arg\min_{z \geq 0} \left(\frac{\gamma}{2}\|Az - b\|_2^2 + \frac{\rho}{2}\|x^{k+1} - z + \frac{\lambda^k}{\rho}\|_2^2 \right).$$

3. **Dual Variable Update**:

$$\lambda^{k+1} := \lambda^k + \rho(x^{k+1} - z^{k+1}).$$

ADMM for Consensus Optimization

The Alternating Direction Method of Multipliers (ADMM) can be generalized to handle *consensus optimization*, where multiple agents or subproblems need to agree on a common solution. The general form of the consensus problem is:

$$\min_{x_1, x_2, \ldots, x_m} \quad \sum_{i=1}^{m} f_i(x_i) \quad \text{s.t.} \quad x_i = z, \ \forall i.$$

Here:

- $f_i(x_i)$: Objective function for the i-th agent or subproblem.
- z: A global consensus variable that all local variables x_i must agree on.

ADMM Steps

To solve this problem, we introduce the consensus variable z and apply ADMM to enforce the constraints $x_i = z$. The augmented Lagrangian is defined as:

$$\mathcal{L}_\rho(x_1, \ldots, x_m, z, \lambda_1, \ldots, \lambda_m) = \sum_{i=1}^{m} \left[f_i(x_i) + \lambda_i^T(x_i - z) + \frac{\rho}{2}\|x_i - z\|_2^2 \right],$$

where:

- λ_i: Dual variable (Lagrange multiplier) for the i-th constraint.
- $\rho > 0$: Penalty parameter for the augmented Lagrangian.

ADMM alternates between updating x_i, z, and λ_i as follows:

1. **Update x_i (local update):**

$$x_i^{k+1} = \arg\min_{x_i} \ f_i(x_i) + \frac{\rho}{2}\|x_i - z^k + \lambda_i^k/\rho\|_2^2.$$

This step solves each $f_i(x_i)$ independently, making it efficient for distributed settings.

2. **Update z (global update):**

$$z^{k+1} = \arg\min_z \sum_{i=1}^m \frac{\rho}{2}\|x_i^{k+1} - z + \lambda_i^k/\rho\|_2^2.$$

   The solution is typically the average:

$$z^{k+1} = \frac{1}{m}\sum_{i=1}^m (x_i^{k+1} + \lambda_i^k/\rho).$$

3. **Update λ_i (dual update):**

$$\lambda_i^{k+1} = \lambda_i^k + \rho(x_i^{k+1} - z^{k+1}).$$

   This adjusts the dual variables to reduce the constraint violation $x_i - z$.

Key Features

- **Parallel Computation:** Each subproblem for x_i can be solved independently and in parallel, making ADMM ideal for distributed systems.
- **Global Consensus:** The z-update ensures that all local variables x_i converge to a single global value z.
- **Flexibility:** ADMM can handle complex objectives $f_i(x_i)$, constraints, or regularizations (e.g., sparsity via ℓ_1-norm).

Extension to Matrix Settings

The consensus ADMM framework extends naturally to matrix settings. The problem is formulated as:

$$\min_{X_1, X_2, \ldots, X_m} \quad \sum_{i=1}^m f_i(X_i) \quad \text{s.t.} \quad X_i = Z, \ \forall i,$$

where:

- $X_i \in \mathbb{R}^{n\times p}$: Matrix-valued variable for the i-th agent or subproblem.
- $f_i(X_i)$: Objective function related to the i-th agent.
- $Z \in \mathbb{R}^{n\times p}$: A global consensus matrix.

Augmented Lagrangian for the Matrix Setting

The augmented Lagrangian for the matrix consensus problem is:

$$\mathcal{L}_\rho(X_1, \ldots, X_m, Z, \Lambda_1, \ldots, \Lambda_m) = \sum_{i=1}^m \left[f_i(X_i) + \mathrm{tr}(\Lambda_i^T(X_i - Z)) + \frac{\rho}{2}\|X_i - Z\|_F^2 \right],$$

where:

- $\Lambda_i \in \mathbb{R}^{n\times p}$: Dual variable (Lagrange multiplier) for the i-th constraint $X_i = Z$.
- $\rho > 0$: Penalty parameter.
- $\|X_i - Z\|_F$: Frobenius norm, defined as $\|A\|_F = \sqrt{\sum_{i,j} A_{i,j}^2}$.
- $\mathrm{tr}(\cdot)$: Trace operator, summing the diagonal elements of a matrix.

ADMM Updates for Matrix Settings

1. **Local Update (X_i):**

$$X_i^{k+1} = \arg\min_{X_i} \ f_i(X_i) + \frac{\rho}{2}\|X_i - Z^k + \Lambda_i^k/\rho\|_F^2.$$

2. **Global Update (Z):**

$$Z^{k+1} = \arg\min_{Z} \sum_{i=1}^{m} \frac{\rho}{2}\|X_i^{k+1} - Z + \Lambda_i^k/\rho\|_F^2.$$

   The solution is:

$$Z^{k+1} = \frac{1}{m}\sum_{i=1}^{m}(X_i^{k+1} + \Lambda_i^k/\rho).$$

3. **Dual Update (Λ_i):**

$$\Lambda_i^{k+1} = \Lambda_i^k + \rho(X_i^{k+1} - Z^{k+1}).$$

Example: Low-Rank Matrix Approximation

Consider approximating a large matrix $M \in \mathbb{R}^{n\times p}$ as $M \approx L$ with a low-rank matrix L. The problem can be formulated as:

$$\min_{L_1,L_2,\ldots,L_m} \sum_{i=1}^{m} \|M_i - L_i\|_F^2 + \lambda\|L_i\|_* \quad \text{s.t.} \quad L_i = Z,\ \forall i,$$

where:

- $\|L_i\|_*$: Nuclear norm (sum of singular values), encouraging low rank.
- M_i: Submatrix of M for agent i.

Using consensus ADMM:

- **Local Update:** Each agent solves its local problem $\|M_i - L_i\|_F^2 + \lambda\|L_i\|_*$ independently.
- **Global Update:** Aggregate the local solutions to update Z.
- **Dual Update:** Adjust the dual variables Λ_i to enforce agreement.

Example: Low-Rank Matrix Approximation

Consider approximating a large matrix $M \in \mathbb{R}^{n\times p}$ as $M \approx L$ with a low-rank matrix L. The problem can be formulated as:

$$\min_{L_1,L_2,\ldots,L_m} \sum_{i=1}^{m} \|M_i - L_i\|_F^2 + \lambda\|L_i\|_* \quad \text{s.t.} \quad L_i = Z,\ \forall i,$$

where:

- $\|L_i\|_*$: Nuclear norm (sum of singular values), encouraging low rank.
- M_i: Submatrix of M for agent i.

Using consensus ADMM:

- **Local Update:** Each agent solves its local problem $\|M_i - L_i\|_F^2 + \lambda\|L_i\|_*$ independently.
- **Global Update:** Aggregate the local solutions to update Z.
- **Dual Update:** Adjust the dual variables Λ_i to enforce agreement.

Discussion on the Penalty Parameter (ρ) in ADMM

The penalty parameter ρ in the Alternating Direction Method of Multipliers (ADMM) plays a crucial role in the convergence and numerical performance of the method. Its selection or adaptation directly impacts the speed, accuracy, and robustness of the algorithm.

1. What Does ρ Do?

In ADMM, the augmented Lagrangian includes the term:

$$\frac{\rho}{2}\|x_i - z\|^2.$$

This term serves two primary purposes:

1. **Penalizes Constraint Violations:** It encourages the primal variables x_i and the consensus variable z to align.
2. **Balances Optimization and Feasibility:** The parameter ρ controls the trade-off between minimizing the objective function $f_i(x_i)$ and enforcing the constraint $x_i = z$.

The choice of ρ affects the convergence behavior:

- **Small ρ:** Emphasizes minimizing the objective function but may allow large constraint violations in early iterations.
- **Large ρ:** Forces constraints to align early, potentially slowing down the optimization of the objective function.

2. Fixed ρ: The Simpler Approach

Using a fixed $\rho > 0$ is the most straightforward approach.

Advantages:

- **Ease of Implementation:** No need to dynamically adjust ρ or monitor residuals.
- **Simpler Analysis:** Convergence proofs and performance bounds are easier to derive.

Challenges:

- **Sensitivity to Problem Properties:**
  - If ρ is too small, the algorithm converges slowly because constraint violations dominate.
  - If ρ is too large, the algorithm may converge poorly or oscillate due to over-enforcing constraints.
- **No Adaptation:** Fixed ρ cannot address problems with inconsistent scales or dynamic requirements.

3. Adaptive ρ: Balancing Convergence

In many practical problems, adapting ρ during iterations improves performance, especially in challenging or ill-conditioned scenarios.

3.1. Residual Balancing Strategy ADMM involves:

- The **primal residual**:
$$r^k = X^k - Z^k,$$
which measures the violation of the constraint $X = Z$.

- The **dual residual**:
$$s^k = \rho(Z^k - Z^{k-1}),$$
which measures the violation of dual feasibility.

To balance these residuals, ρ can be updated as:

$$\rho^{k+1} = \begin{cases} \tau\rho^k, & \text{if } \|r^k\| > \mu\|s^k\|, \\ \rho^k/\tau, & \text{if } \|s^k\| > \mu\|r^k\|, \\ \rho^k, & \text{otherwise,} \end{cases}$$

where $\tau > 1$ and $\mu > 1$ are user-defined constants.

3.2. Incremental Increase Another simple strategy is to gradually increase ρ:

$$\rho^{k+1} = \rho^k(1 + \epsilon),$$

where $\epsilon > 0$ is a small constant.

3.3. Heuristic Adaptation Heuristic strategies can adjust ρ dynamically based on problem properties, such as condition numbers or scaling. For example:

- Start with a small ρ, and increase it when residuals stagnate.
- Use scaling factors based on the problem dimensions.

4. Theoretical Considerations

- **Fixed ρ:** ADMM is guaranteed to converge for convex problems under mild conditions if ρ is fixed. However, the convergence rate is highly sensitive to the choice of ρ.
- **Adaptive ρ:** Theoretical convergence guarantees are achievable under specific rules for adjusting ρ (e.g., bounded changes, non-decreasing sequences).
- **Non-Convex Problems:** In non-convex problems, ρ still affects convergence, but guarantees are weaker. Adaptive ρ becomes more critical in such cases.

5. Practical Recommendations

1. **Start with a Fixed ρ:** Choose a moderately tuned ρ based on problem properties. A common heuristic is to scale ρ based on the magnitudes of the objective and constraints.
2. **Switch to Adaptive ρ:** Monitor the primal and dual residuals during optimization. If residuals are imbalanced or convergence stagnates, switch to an adaptive strategy.
3. **Tune Adaptive Parameters:** For residual balancing, typical values are $\tau = 2$ and $\mu = 10$. Ensure ρ does not change too drastically between iterations.
4. **Scale ρ with Problem Size:** For large-scale problems, consider scaling ρ based on problem dimensions or the norm of the data.

6. Example: Matrix Consensus

Consider the matrix consensus problem:

$$\min_{X_1, X_2, \ldots, X_m} \sum_{i=1}^{m} f_i(X_i) \quad \text{s.t.} \quad X_i = Z, \ \forall i.$$

- **Fixed ρ:** A fixed ρ enforces the constraints $X_i = Z$ uniformly across iterations. Poorly chosen ρ can slow convergence.
- **Adaptive ρ:** Early iterations prioritize minimizing $f_i(X_i)$ with smaller ρ, while later iterations enforce consensus $X_i = Z$ with larger ρ.

7. Conclusion

The penalty parameter ρ in ADMM is critical for balancing convergence speed and numerical stability. While fixed ρ is simple and effective for well-conditioned problems, adaptive strategies offer significant advantages in challenging cases. Monitoring residuals and dynamically adjusting ρ enables faster and more robust convergence.