

12.1 The Forward Process

Jump to x_t in one shot. $\bar{\alpha}$ from 1 to 0. The kernel is not learned.

These notes match the lecture slides. Use **(slides)** on the course hub for the deck.

A **diffusion** model learns to denoise. The forward process is not learned. You add Gaussian noise on a schedule until the sample is almost prior noise. Training later asks a net to predict that noise.

First time this method appears. Diffusion (DDPM) **destroys** a sample with Gaussian noise on a schedule, then later learns to reverse it.

What. Forward process $q(x_t | x_0)$ is a **fixed** Gaussian that interpolates from data ($t = 0$) to nearly isotropic noise ($t = T$). You can **jump** to any t in one shot. **Why.** GANs fight a min-max. Diffusion trains a denoiser with MSE. The kernel is not learned. **Architecture.** Markov chain

$x_0 \rightarrow x_1 \rightarrow \dots \rightarrow x_T$. Variance schedule β_t or $\bar{\alpha}_t \in [1, 0]$. **How.** Sample $\varepsilon \sim \mathcal{N}(0, I)$, set $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\varepsilon$. Lab 12 does this in 2-D. **Formula.** $q(x_t | x_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$.

Tradeoffs. + Stable training target, likelihood-related. – Forward is easy; reverse sampling is slow (next notes). Not a GAN.

1. Destroy the sample on purpose

Start from data x_0 . A variance schedule $0 < \beta_t < 1$ (small at first, larger later) defines


$$\alpha_t = 1 - \beta_t, \quad \bar{\alpha}_t = \prod_{s=1}^t \alpha_s.$$

The Markov step is $q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t} x_{t-1}, \beta_t I)$. You do not have to simulate every hop. The marginal is closed form:

$$q(x_t | x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t)I).$$

In code: $x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \varepsilon$ with $\varepsilon \sim \mathcal{N}(0, I)$.

Forward process: data becomes noise



That one-liner is the whole point of this note. Training (note **12.2**) samples a random t , draws ε , forms x_t in one shot, and asks a net to guess ε . You never need to walk $1, 2, \dots, t$ unless you want a picture of

the path.

2. What the schedule is doing

$\bar{\alpha}_t$ starts near 1 (the point is still itself) and falls toward 0 (the point is almost ε). Linear, cosine, and learned β_t are all ways to spend more steps where structure dies. Too fast a schedule and early t is already junk; too slow and you waste capacity on near-copies of x_0 .

T (often 1000 in image papers) is a discretization. Lab 12 uses a small T in 2-D so you can plot the path. Print $\bar{\alpha}_1, \bar{\alpha}_{T/2}, \bar{\alpha}_T$: they should fall.

3. Why this is nicer than a GAN to train

The forward kernel is **fixed**. There is no discriminator to race. The training signal at each t is a denoising regression against a known ε . You still have to pick $\{\beta_t\}$ and T . Sampling will be a reverse chain (note **12.2**), which is slower than one GAN forward pass.

Coverage is not automatic, but you are not playing a min-max game. That is the trade this week is for.

4. Teaching this note

About **35 minutes** at the board, then **~6 minutes** of video. First of three Week-12 notes.

- **0–12 min.** Markov step vs closed-form marginal. Write α_t and $\bar{\alpha}_t$.
 - **12–22 min.** The one-line reparameterization. Draw x_0 mixing with ε .
 - **22–33 min.** Worked scalar step. Change $\bar{\alpha}_t$ from 1 to 0 and watch x_t .
 - **Then** play Umar Jamil **7:35–13:00** (forward/reverse cartoon, then the math of q). Pause on the closed-form $q(x_t | x_0)$.
-

5. Worked example

Take a 1-D point $x_0 = 1.0$ and a draw $\varepsilon = 2.0$. Then

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \varepsilon.$$

$\bar{\alpha}_t$	$\sqrt{\bar{\alpha}_t} x_0$	$\sqrt{1 - \bar{\alpha}_t} \varepsilon$	x_t
1	1.0	0	1.0
0.64	0.8	$0.6 \times 2.0 = 1.2$	2.0
0	0	2.0	2.0

At $\bar{\alpha}_t = 0.64$ you already sit at 2.0: more noise than signal. At $\bar{\alpha}_t = 0$ you have thrown x_0 away; x_t is just ε .

Same arithmetic in Lab 12 with a vector $x_0 = (1.5, -0.8)$ and a vector ε . The square roots are scalar; they multiply every coordinate.

You can jump to any t without simulating $q(x_s | x_{s-1})$ for $s = 1, \dots, t$. That is why training is a single Gaussian draw, not a loop of length t .

 Three $\bar{\alpha}_t$ values on $x_0 = 1, \varepsilon = 2$

 src="graphics/12.1-diffusion-forward/alphabar.png" />

Pick a noise level, corrupt x , then learn to undo a bit. The closed-form jump is why the loss never walks $1, \dots, t$.

6. Where students get stuck

- Simulating all T Markov steps at train time. You sample t and use $q(x_t | x_0)$.
- Mixing up β_t , α_t , and $\bar{\alpha}_t$. Write the three definitions every time.
- Setting every β_t huge. Then $\bar{\alpha}_t$ crashes in a few steps and the net never sees a lightly noised x_0 .

7. Video

Watch [Umar Jamil, How diffusion models work](#), 7:35–13:00.

Pause on the forward process q (no θ) and the closed-form jump to x_t . Skip the U-Net coding (after ~16:30) in this class. Note **12.2** continues the same video at the training/sampling chapters.

8. Practice

1. If $\bar{\alpha}_t = 1$, what is x_t ? If $\bar{\alpha}_t = 0$, what is x_t ?
2. Why can you sample x_t in one line instead of looping t steps of $q(x_s | x_{s-1})$?
3. A classmate sets every $\beta_t = 0.9$. What goes wrong in the first few t ?
4. $x_0 = 4$, $\bar{\alpha}_t = 0.25$, $\varepsilon = -1$. Compute x_t . ($\sqrt{0.25} = 0.5$, $\sqrt{0.75} \approx 0.866$.)
5. $\beta_1 = 0.1$, $\beta_2 = 0.2$. Compute $\alpha_1, \alpha_2, \bar{\alpha}_2$.