

1.4 Word Embeddings and Semantic Geometry

Dense vectors from a corpus window, skip-gram softmax, analogies, and a signed-projection bias probe.



These notes match the lecture slides. Use **(slides)** on the course hub for the deck.

A language model does not read characters as characters. It reads **vectors**. This week those vectors are word embeddings. From Week 3 they are token embeddings inside a transformer. The geometry is the same idea.

1. One-hot is a bad geometry

Give each word its own coordinate. *film* and *movie* are then orthogonal: inner product 0. The vector is sparse, as long as the vocabulary, and a new word has no coordinate at all.

A **distributed** representation is short and dense (50–300 dimensions is a classical size; transformer hidden sizes are larger). Meaning is spread across coordinates. Similar usage → similar vectors.



Distributional is about the data: words that occur in similar contexts mean similar things (Harris).

Distributed is about the code: a dense vector. An **embedding** is the map from a token id into that vector.

Write those three words on the board as a triangle. Mixing them is the most common exam slip this week.

2. How a distributional embedding is obtained


You do not type a definition. You **count or predict company**.

1. Take a corpus of running text.

2. At each position, treat one word as the **center** and the nearby words as **context** (a window).
3. Train a fake task: from the center, predict a neighbor (skip-gram), or the reverse (CBOW). Words that appear in similar windows get similar vectors.
4. After training, **keep the center rows**. Those rows are the embeddings. Throw away the softmax classifier.

That is why *cat* sits near *dog*: they kept similar neighbors, not because someone labeled them as animals.

Distributional embeddings are obtained from running text, not from a thesaurus




Firth: you shall know a word by the company it keeps. Word2Vec (next section) is one algorithm for that. Counting co-occurrences and then factoring the matrix is another. Both are distributional. This course uses the predictive story because it is the same training loop as an LLM.

3. Semantic geometry

If the map is good, **direction means something**. The Mikolov analogy is the slogan:

$$\overrightarrow{\text{king}} - \overrightarrow{\text{man}} + \overrightarrow{\text{woman}} \approx \overrightarrow{\text{queen}}.$$

The same offset can look like gender, tense, or capital-of. You measure similarity with **cosine**, not Euclidean distance (length is partly frequency).



Nearest neighbors are a debug tool. If *Paris* sits next to *France* and *Rome*, training learned something. If it sits next to punctuation, the corpus or the window is wrong.


Cosine of u, v is $u^T v / (\|u\| \|v\|)$. It ignores length. That is why a frequent word with a long vector can still sit next to a rare synonym.

4. How Word2Vec learns the map

You do not build a huge co-occurrence matrix and then factor it (that is LSA / GloVe; we do not assign SVD this week). **Word2Vec** trains a small two-layer net on a fake task.

Skip-gram: from the center word, predict each neighbor in a window. **CBOW**: from the neighbors, predict the center.

Skip-gram: predict each neighbor from the center word




Let v_c be the **input** (center) embedding of the center word, and u_o the **output** embedding of a neighbor. Skip-gram's softmax is Mikolov §§1–3:

$$P(w_o | w_c) = \frac{\exp(u_o^\top v_c)}{\sum_w \exp(u_w^\top v_c)}.$$

That is the same softmax as note 1.2, with logits $z_w = u_w^\top v_c$.

Tiny vocabulary {cat, mat, sat}. Take center *the* with $v_c = (1, 0)$, and output vectors $u_{\text{cat}} = (1, 0)$, $u_{\text{mat}} = (0.5, 0)$, $u_{\text{sat}} = (0, 1)$. The three scores are 1, 0.5, 0. Softmax:

$$p \propto (e^1, e^{0.5}, e^0) \approx (2.718, 1.649, 1), \quad p \approx (0.51, 0.31, 0.19).$$



Input is a token id. `nn.Embedding(V, d)` is a lookup table: row c of a $V \times d$ matrix. A one-hot times that matrix is the same map. After training, Word2Vec **throws away the softmax** (U) and keeps the center rows v_w . Those rows are the embeddings you plot.

A full softmax over 100,000 words is slow. Negative sampling (Mikolov 2013b) replaces it. We cite that paper; we do not derive it this week.

A transformer does not throw the rest of the net away: it *keeps* transforming those vectors. Word2Vec is the geometry without the rest of the stack.


Counting vs. predicting. Classical vector semantics (Jurafsky Ch. 5) can start from a co-occurrence matrix. Skip-gram *predicts* context instead of counting it. Both are distributional. This course uses the predictive story because it is the same training loop as an LLM.

5. Intrinsic versus extrinsic evaluation

Intrinsic: does the space look right on its own? Analogies, word similarity datasets, clustering of synsets.

Extrinsic: plug the vectors into a downstream model (sentiment, retrieval, NER). If the task score moves, the embedding helped that task.

Two ways to score embeddings. They need not agree.




They need not agree. A space that wins analogies can still be a weak feature for your classifier. For a project, extrinsic is the number that matters. Intrinsic is how you debug.

6. Bias is also geometry

The same offsets that encode "capital of" can encode stereotypes. Occupation words often sit along a she/he direction in older embeddings. That is not a bug in cosine; it is a property of the training text.

Geometry can encode stereotypes as well as meaning



Lab 1's probe is the signed projection onto $o = \vec{he} - \vec{she}$:

$$\text{score}(v) = \frac{v^T o}{\|o\|}.$$

A larger score sits closer to *he* in this constructed space. That is not a bug in cosine; it is a property of the (here, invented) geometry. Debiasing is incomplete; reporting the probe is part of the project

writeup when your system ranks people, jobs, or medical outcomes.

7. Teaching this note

~18 minutes. One-hot vs dense, then the four-box pipeline (corpus → window → fake task → keep the rows), then skip-gram's $P(o | c)$ on the $V = 3$ softmax, then cosine, then the occupation projection so Lab 1's probe is not a surprise. StatQuest Word2Vec is **homework (0:00–12:00, skip-gram vs. CBOW)**. Do not play CS224N Lecture 2 in class.

8. Worked example

Three 2-D vectors:

$$\overrightarrow{\text{movie}} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \quad \overrightarrow{\text{film}} = \begin{bmatrix} 0.8 \\ 0.2 \end{bmatrix}, \quad \overrightarrow{\text{pizza}} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}.$$

$\text{Cosine}(\text{movie}, \text{film}) = 0.8 / \sqrt{0.68} \approx 0.97$. $\text{Cosine}(\text{movie}, \text{pizza}) = 0$. One-hot versions of the same three words would all be orthogonal.

Bias probe aligned with Lab 1: $o = \overrightarrow{\text{he}} - \overrightarrow{\text{she}}$. If $\overrightarrow{\text{he}} = (1, 0)$, $\overrightarrow{\text{she}} = (0, 1)$, then $o = (1, -1)$ and $\|o\| = \sqrt{2}$. For $\overrightarrow{\text{engineer}} = (0.9, 0.1)$ and $\overrightarrow{\text{nurse}} = (0.1, 0.8)$,

$$\text{score}(\text{engineer}) = \frac{0.9 - 0.1}{\sqrt{2}} \approx 0.57, \quad \text{score}(\text{nurse}) = \frac{0.1 - 0.8}{\sqrt{2}} \approx -0.49.$$

Engineer sits toward *he*; nurse sits toward *she*. The CSV in Lab 1 was built to make that split visible. On a real model you report the same probe, not this cartoon.

9. Where students get stuck

- Using Euclidean distance and then “the frequent word is never nearest.”
 - Calling the one-hot vector an embedding because it is a vector.
 - Treating analogy accuracy as the project metric.
 - Thinking an embedding file appeared without a corpus: skip the four-box pipeline.
-

10. Video

Watch [StatQuest: Word Embedding and Word2Vec, Clearly Explained](#).

Pause on skip-gram vs. CBOW, and on the moment the hidden weights become the embedding table. You can skip the neural-net-from-scratch recap if note 1.2 already landed.

11. Practice

1. In one sentence each: distributional, distributed, embedding.
2. Why is cosine the default similarity?
3. Give one intrinsic and one extrinsic test for a word embedding. Which one would you put in a project report?
4. Compute cosine between $a = \begin{bmatrix} 3 \\ 4 \end{bmatrix}$ and $b = \begin{bmatrix} 4 \\ 3 \end{bmatrix}$. Then compute $a^\top b$. Which number changed more when you replace b by $10b$?
5. In 2-D, $\overrightarrow{\text{king}} = \begin{bmatrix} 2 \\ 2 \end{bmatrix}$, $\overrightarrow{\text{man}} = \begin{bmatrix} 2 \\ 0 \end{bmatrix}$, $\overrightarrow{\text{woman}} = \begin{bmatrix} 0 \\ 2 \end{bmatrix}$. Compute $\overrightarrow{\text{king}} - \overrightarrow{\text{man}} + \overrightarrow{\text{woman}}$. What point would you hope $\overrightarrow{\text{queen}}$ near?
6. With $v_c = (1, 0)$ and scores $u^\top v_c = (1, 0.5, 0)$ for cat, mat, sat, which word does skip-gram call most likely? Write the three softmax probabilities to two decimals.
7. In four short lines: how is a distributional embedding obtained from a corpus? (Corpus, window, fake task, what you keep.)