

1.3 Gradient Descent

Forward, loss, backward, update, and one ReLU unit's chain rule. That loop is pretraining.

These notes match the lecture slides. Use **(slides)** on the course hub for the deck.

Training a net is an optimization problem: pick weights so a **loss** is small on training batches, then check a held-out set. Gradient descent is the default way to move the weights.

First time this method appears. Gradient descent is how nets learn. Answer these before the ReLU chain rule.

What. Repeatedly walk downhill on a loss $L(\theta)$. Stochastic GD uses a minibatch, not the full dataset.

Why. You cannot solve $\nabla L = 0$ in closed form for a deep net. This is the training loop of pretraining.


Architecture. Forward (compute L) → backward (compute $\nabla_{\theta} L$) → update θ . **How.** Pick a step size η .

Too large: diverge. Too small: crawl. Adam is a later variant with per-coordinate scales. **Formula.**

$\theta \leftarrow \theta - \eta \nabla_{\theta} L$. Chain rule on one ReLU: $\partial L / \partial w = (\partial L / \partial a) 1_{z > 0} x$. **Tradeoffs.** + Simple, scales to millions of steps. – Local minima / saddles, needs a step-size choice, noisy minibatches.

1. Loss as a surface

Collect all weights and biases into one vector θ . A loss $L(\theta)$ is a scalar. For classification it is often **cross-entropy**; for a regression toy it is squared error. In two dimensions you can draw the contours. Gradient descent follows $-\nabla L$.



The picture is a cartoon. Real nets live in millions of dimensions. The geometry is still the same idea: the gradient is the direction of steepest increase, so you walk the other way.

Cross-entropy on a next-token distribution is the LLM default. For one token with predicted probability p on the correct class,

$$L = -\log p.$$

Lab 1 uses binary cross-entropy on **logits** z , not on $\sigma(z)$. PyTorch's `BCEWithLogitsLoss` is

$$L = \max(z, 0) - zy + \log(1 + e^{-|z|}),$$

which is numerically stable BCE. The extra line you need: **pass z , never $\text{sigmoid}(z)$, into that loss.** If you sigmoid first, you squash twice.


You need: **one scalar**, then **one vector of derivatives**. Nielsen calls this the **cost**; the usual name here is **loss**; a statistician calls it **empirical risk**. Same object.

2. The update

One step is

$$\theta \leftarrow \theta - \eta \nabla_{\theta} L(\theta).$$

η is the **learning rate**. Too small: you crawl. Too large: you jump over the valley or diverge.



Stochastic gradient descent (SGD) estimates ∇L on a **minibatch**, not the full dataset. That is how you train on Wikipedia-scale text. Adam and related methods rescale coordinates using a running average of gradients; you will use them in PyTorch without deriving them today.


A noisy minibatch gradient is not a bug. It is the only gradient you can afford, and the noise can help you leave sharp spikes. The price is that one step is not guaranteed to decrease L .

3. Forward, loss, backward, update

One training iteration is four moves:

1. **Forward.** Compute activations layer by layer. Keep them; the backward pass needs them.
2. **Loss.** Compare $\hat{y}$ to the target y .
3. **Backward.** Chain rule from the loss to every weight (**backpropagation**).
4. **Update.** Take a descent step.

Training a net is this loop. Backprop is just the chain rule.



PyTorch does (3) for you if the forward pass used `nn.Module` and a differentiable loss. Your job is to write (1), pick L , and call `optimizer.step()`.

If you forget `zero_grad()`, gradients **accumulate** across steps. That is a silent η disaster, not a new algorithm.

4. One ReLU unit, by hand

Treat each operation as a **gate** with a local derivative. Autograd is that picture. For squared error on one ReLU neuron,


$$z = w^T x + b, \quad a = \text{ReLU}(z), \quad L = \frac{1}{2}(a - y)^2.$$

Chain rule, using $\text{ReLU}'(z) = 1_{z > 0}$ from note 1.2:

$$\frac{\partial L}{\partial w} = (a - y) \text{ReLU}'(z) x.$$

Same pattern for b , with x replaced by 1. If $z \leq 0$, the local derivative is 0 and that example does not move w . That is a dead ReLU on this point, not a PyTorch bug.

One ReLU unit: cache x and z ; autograd is this chain



backward: $(a - y) \cdot \mathbf{1}_{z > 0} \cdot x$ (CS231N: local $\times$ upstream)

The softmax/cross-entropy cousin is $\partial L / \partial w = (p - y)x$. Same "upstream error times input" shape. We do not need Adam's derivation this week.

5. What this has to do with language models

Pretraining an LLM is this loop on a next-token loss, at huge batch and data scale. Fine-tuning (Week 8) is the same loop on a smaller, more specific dataset. RLHF (Week 9) changes what L is, not the fact that you are differentiating a scalar with respect to weights.

If the gradient is noise, or η is wrong, or the net is linear, no amount of later "prompt engineering" will save the representation. Get this loop right on a tiny problem in Lab 1 before you touch Hugging Face.

6. Teaching this note

~18 minutes. Draw a 1-D parabola, write the update, do **two** numeric steps plus the overshoot, then the four-box loop and `zero_grad`. Spend four minutes on the one-ReLU chain rule and the `BCEWithLogitsLoss` warning. 3Blue1Brown gradient descent (**0:00–12:00**) and the backprop follow-up are **homework**. Do not derive a softmax Jacobian in this block.

7. Worked example

Let $L(\theta) = (\theta - 3)^2$, start at $\theta_0 = 0$, take $\eta = 0.25$. Then $\nabla L = 2(\theta - 3)$.

Step 1:

$$\theta_1 = 0 - 0.25 \cdot 2(0 - 3) = 0 - 0.25 \cdot (-6) = 1.5.$$

Step 2:

$$\theta_2 = 1.5 - 0.25 \cdot 2(1.5 - 3) = 1.5 - 0.5 \cdot (-1.5) = 2.25.$$

The minimum is at 3. You moved $0 \rightarrow 1.5 \rightarrow 2.25$. If you instead take $\eta = 2$, step 1 is $0 - 2 \cdot (-6) = 12$, which **overshoots**. Draw both arrows on the same parabola.

One-unit check: $x = 2$, $w = 0.5$, $b = 0$, $y = 1$. Then $z = 1$, $a = \text{ReLU}(1) = 1$, $L = 0$, so $\partial L / \partial w = 0$. Change y to 0: $a - y = 1$, $\text{ReLU}'(1) = 1$, $\partial L / \partial w = 1 \cdot 1 \cdot 2 = 2$. One SGD step with $\eta = 0.1$ sends $w \leftarrow 0.5 - 0.1 \cdot 2 = 0.3$.

8. Where students get stuck

- Walking **up** the gradient (forgetting the minus sign).
 - Thinking SGD is “wrong GD” rather than the scalable estimator of ∇L .
 - Dropping activations after the forward pass, then being surprised that backward needs them.
 - Passing `sigmoid(z)` into `BCEWithLogitsLoss` (the loss already includes the sigmoid).
-

9. Video

Watch [3Blue1Brown: Gradient descent, how neural networks learn](#).

Pause when the ball follows $-\nabla L$, and when a large step jumps the valley. Optional: [3Blue1Brown: What is backpropagation really doing?](#) after class (~14 min); pause on the chain-rule diagram, not the code.

10. Practice

1. You double η and the loss oscillates. What happened geometrically?
2. Why keep $a^{(1)}$ after you already have $\hat{y}$?
3. SGD uses a minibatch. Name one reason that is not “the full gradient is too slow.”
4. For $L(\theta) = \theta^2$, $\theta_0 = 4$, $\eta = 0.1$, write θ_1 and θ_2 . (Use $\nabla L = 2\theta$.)
5. Same L and $\theta_0 = 4$, but $\eta = 1.1$. Compute θ_1 . Did the step move closer to 0 or farther?
6. For the one-ReLU unit with $x = 2$, $w = 0.5$, $b = 0$, $y = 0$, write $\partial L / \partial w$. Then take $\eta = 0.1$ and the new w .