

1.3 Gradient Descent


Forward, loss, backward, update, and one ReLU unit's chain rule. That loop is pretraining.

These notes match the lecture slides. Use **(slides)** on the course hub for the deck.

Training a net is an optimization problem: pick weights so a **loss** is small on training batches, then check a held-out set. Gradient descent is the default way to move the weights.

1. Loss as a surface

Collect all weights and biases into one vector θ . A loss $L(\theta)$ is a scalar. For classification it is often **cross-entropy**; for a regression toy it is squared error. In two dimensions you can draw the contours. Gradient descent follows $-\nabla L$.



The picture is a cartoon. Real nets live in millions of dimensions. The geometry is still the same idea: the gradient is the direction of steepest increase, so you walk the other way.

Cross-entropy on a next-token distribution is the LLM default. For one token with predicted probability p on the correct class,

$$L = -\log p.$$

Lab 1 uses binary cross-entropy on **logits** z , not on $\sigma(z)$. PyTorch's `BCEWithLogitsLoss` is

$$L = \max(z, 0) - zy + \log(1 + e^{-|z|}),$$

which is numerically stable BCE. The extra line you need: **pass z , never `sigmoid(z)`, into that loss.** If you sigmoid first, you squash twice.


You need: **one scalar**, then **one vector of derivatives**. Nielsen calls this the **cost**; CS231N calls it the **loss**; a statistician calls it **empirical risk**. Same object.

2. The update

One step is

$$\theta \leftarrow \theta - \eta \nabla_{\theta} L(\theta).$$

η is the **learning rate**. Too small: you crawl. Too large: you jump over the valley or diverge.



Stochastic gradient descent (SGD) estimates ∇L on a **minibatch**, not the full dataset. That is how you train on Wikipedia-scale text. Adam and related methods rescale coordinates using a running average of gradients; you will use them in PyTorch without deriving them today.


A noisy minibatch gradient is not a bug. It is the only gradient you can afford, and the noise can help you leave sharp spikes. The price is that one step is not guaranteed to decrease L .

3. Forward, loss, backward, update

One training iteration is four moves:

1. **Forward.** Compute activations layer by layer. Keep them; the backward pass needs them.
2. **Loss.** Compare $\hat{y}$ to the target y .
3. **Backward.** Chain rule from the loss to every weight (**backpropagation**).
4. **Update.** Take a descent step.

Training a net is this loop. Backprop is just the chain rule.



PyTorch does (3) for you if the forward pass used `nn.Module` and a differentiable loss. Your job is to write (1), pick L , and call `optimizer.step()`.

If you forget `zero_grad()`, gradients **accumulate** across steps. That is a silent η disaster, not a new algorithm.

4. One ReLU unit, by hand

CS231N treats each operation as a **gate** with a local derivative. Autograd is that picture. For squared error on one ReLU neuron,


$$z = w^T x + b, \quad a = \text{ReLU}(z), \quad L = \frac{1}{2}(a - y)^2.$$

Chain rule, using $\text{ReLU}'(z) = 1_{z > 0}$ from note 1.2:

$$\frac{\partial L}{\partial w} = (a - y) \text{ReLU}'(z) x.$$

Same pattern for b , with x replaced by 1. If $z \leq 0$, the local derivative is 0 and that example does not move w . That is a dead ReLU on this point, not a PyTorch bug.

One ReLU unit: cache x and z ; autograd is this chain



backward: $(a - y) \cdot 1_{z > 0} \cdot x$ (CS231N: local $\times$ upstream)

CMU 11-711 writes the softmax/cross-entropy cousin $\partial L / \partial w = (p - y)x$. Same "upstream error times input" shape. We do not need Adam's derivation this week.

5. What this has to do with language models

Pretraining an LLM is this loop on a next-token loss, at huge batch and data scale. Fine-tuning (Week 8) is the same loop on a smaller, more specific dataset. RLHF (Week 9) changes what L is, not the fact that you are differentiating a scalar with respect to weights.

If the gradient is noise, or η is wrong, or the net is linear, no amount of later “prompt engineering” will save the representation. Get this loop right on a tiny problem in Lab 1 before you touch Hugging Face.

6. Teaching this note

~**18 minutes**. Draw a 1-D parabola, write the update, do **two** numeric steps plus the overshoot, then the four-box loop and `zero_grad`. Spend four minutes on the one-ReLU chain rule and the `BCEWithLogitsLoss` warning. 3Blue1Brown gradient descent (**0:00–12:00**) and the backprop follow-up are **homework**. Do not derive a softmax Jacobian in this block.

7. Worked example

Let $L(\theta) = (\theta - 3)^2$, start at $\theta_0 = 0$, take $\eta = 0.25$. Then $\nabla L = 2(\theta - 3)$.

Step 1:

$$\theta_1 = 0 - 0.25 \cdot 2(0 - 3) = 0 - 0.25 \cdot (-6) = 1.5.$$

Step 2:

$$\theta_2 = 1.5 - 0.25 \cdot 2(1.5 - 3) = 1.5 - 0.5 \cdot (-1.5) = 2.25.$$

The minimum is at 3. You moved $0 \rightarrow 1.5 \rightarrow 2.25$. If you instead take $\eta = 2$, step 1 is $0 - 2 \cdot (-6) = 12$, which **overshoots**. Draw both arrows on the same parabola.

One-unit check: $x = 2$, $w = 0.5$, $b = 0$, $y = 1$. Then $z = 1$, $a = \text{ReLU}(1) = 1$, $L = 0$, so $\partial L / \partial w = 0$. Change y to 0: $a - y = 1$, $\text{ReLU}'(1) = 1$, $\partial L / \partial w = 1 \cdot 1 \cdot 2 = 2$. One SGD step with $\eta = 0.1$ sends $w \leftarrow 0.5 - 0.1 \cdot 2 = 0.3$.

8. Where students get stuck

- Walking **up** the gradient (forgetting the minus sign).
 - Thinking SGD is “wrong GD” rather than the scalable estimator of ∇L .
 - Dropping activations after the forward pass, then being surprised that backward needs them.
 - Passing `sigmoid(z)` into `BCEWithLogitsLoss` (the loss already includes the sigmoid).
-

9. Video

Watch [3Blue1Brown: Gradient descent, how neural networks learn](#).

Pause when the ball follows $-\nabla L$, and when a large step jumps the valley. Optional: [3Blue1Brown: What is backpropagation really doing?](#) after class (~14 min); pause on the chain-rule diagram, not the code.

10. Practice

1. You double η and the loss oscillates. What happened geometrically?
2. Why keep $a^{(1)}$ after you already have $\hat{y}$?
3. SGD uses a minibatch. Name one reason that is not "the full gradient is too slow."
4. For $L(\theta) = \theta^2$, $\theta_0 = 4$, $\eta = 0.1$, write θ_1 and θ_2 . (Use $\nabla L = 2\theta$.)
5. Same L and $\theta_0 = 4$, but $\eta = 1.1$. Compute θ_1 . Did the step move closer to 0 or farther?
6. For the one-ReLU unit with $x = 2$, $w = 0.5$, $b = 0$, $y = 0$, write $\partial L / \partial w$. Then take $\eta = 0.1$ and the new w .