

14.1 Retrieval-Augmented Generation

First time: what RAG is, why, architecture, cosine and Lewis's sum, then the three-doc ranking.

These notes match the lecture slides. Use **(slides)** on the course hub for the deck.

This is the first full lecture on **retrieval-augmented generation (RAG)**. Treat the method as new. Week **10.3** only named it so RAFT would make sense. Today you should be able to say what RAG is, why anyone bothers, how the boxes connect, which formulas you write, and what it is good and bad at.

1. What RAG is

A frozen language model answers from **weights**. Those weights were trained on a snapshot of the web. They do not contain this semester's syllabus, a private PDF, or a fact that changed last week. When the model is unsure, it still produces fluent text. That fluent guess is a **hallucination**.

RAG means: **look up relevant text first, then generate the answer from that text.**

The name is literal.

- **Retrieval:** given a question, find a few passages from a collection you control.
- **Augmented:** put those passages into the prompt, next to the question.
- **Generation:** the same next-token model as always writes the answer, now **conditioned** on the passages.

You are not training a new network today. You are changing what the network is allowed to read at request time. Stanford **CS224N 2026 L10** calls the passages **non-parametric memory**: the facts live in an index you can edit, not only inside θ .

A search engine returns links. RAG returns an **answer**, with the retrieved passages sitting in the context window. A fine-tune changes weights. RAG leaves the weights alone and changes the **context**.

2. Why we use it

Three jobs keep coming up in products and in this course.

1. **Fresh or private knowledge.** Campus hours, a company policy PDF, a paper that appeared this month. Retraining a 7B model every time a PDF moves is the wrong tool.
2. **Citations.** If the answer is supposed to come from chunk [d1], a reader can open [d1] and check. Weights do not give you that handle.
3. **A smaller generator plus a large library.** You can keep a modest LLM and a large corpus, instead of stuffing the corpus into parameters.

Why not just fine-tune on the PDFs? Fine-tuning is slow to update, easy to overfit, and still does not cite a chunk id. Use it for *style* and *format* (Week 8 SFT / LoRA). Use RAG for *facts that live in documents*.

Why not dump the whole corpus into a long prompt? Context is expensive, and models often ignore the middle of a long list (**lost in the middle**). Retrieval is how you pick a handful of passages instead of all of them.

Karpathy's browser demo is the same idea with the web as the index: emit a search, read the hits, then write. Custom GPTs that file your PDFs are the same loop with a private index.

3. Architecture

RAG is a **pipeline**, not a new layer inside the transformer.

Offline (once, then whenever the corpus changes):


1. Collect documents (HTML, PDF, notes, a CSV of snippets).
2. **Chunk** them into passages of a few hundred tokens, often with overlap so a sentence is not split from its date.
3. **Embed** each chunk with the same embedder you will use for queries (Week 1 vectors; Week 5 CLIP is the multimodal cousin). Lab 14 uses bag-of-words counts so you can print the ranking without a GPU.
4. Store, for each chunk: the **vector**, the **raw text**, and an **id**. That triple is the **index**. Vectors without text cannot be cited.

Online (every question):

1. Embed the query with the **same** embedder.
2. Score every chunk (or a fast approximate neighbor search). Keep the top- k .
3. **Stuff** the k texts into a prompt: question, passages with ids, and instructions (answer from the passages; cite ids; say when the context is missing).
4. The language model generates tokens. It is still $p(y_t \mid y_{<t}, \text{prompt})$. RAG changed the prompt, not the architecture.

RAG architecture: the transformer is unchanged; the prompt is not

Offline (build the index)



Online (every question)



Dense embeddings are one index. **Keyword overlap / BM25** is another (Lab 10's toy retriever). Hybrid systems use both. The generator does not care which ranker you used; it only sees the strings you stuffed.

4. How it works, step by step

Take a question students will actually type: *"What time does Bender Library close on a weekday?"*

1. **Index already exists.** Suppose chunk [d1] says the weekday close is 23:00, [d3] says something about hours but not 23:00, [d2] is about a shuttle.
2. **Embed the question.** In Lab 14 this is a count vector over the vocabulary. In a product this is a neural embedding.
3. **Rank.** Cosine (below) puts [d1] first, [d3] second, [d2] last.
4. **Choose k .** $k = 1$ stuffs only [d1]. $k = 2$ also stuffs [d3]. Recall can rise; a **near-miss** can now mislead the generator.
5. **Prompt.** Something like: *Use only the passages. Cite ids like [d1]. If the passages do not contain the answer, say you do not know.* Then list the passages.
6. **Generate.** A good run answers 23:00 and cites [d1]. A bad run cites [d2] while using the 23:00 fact, or ignores [d1] and invents 22:00 from pretraining.

The generator is still a next-token model. A wrong index cannot be saved by a prettier prompt.

5. Mathematical formulas

Let $\mathcal{Z}$ be the set of chunks. Let $e(\cdot)$ be the embedder (counts or a neural map). For a query q and a chunk z ,

$$\cos(e(q), e(z)) = \frac{e(q)^\top e(z)}{\|e(q)\| \|e(z)\|}.$$

Retrieve the k chunks with largest cosine (ties: pick a rule and stick to it):

$$\hat{z}_{1:k} = \text{TopK}_{z \in \mathcal{Z}} \cos(e(q), e(z)).$$

Generate one token at a time from the stuffed prompt $x = \text{Prompt}(q, \hat{z}_{1:k})$:

$$p_\theta(y \mid q, \hat{z}_{1:k}) = \prod_t p_\theta(y_t \mid y_{<t}, x).$$

That is what Lab 14 and most production “stuffing” RAG systems do: **one** generation given the concatenated top- k .

Lewis et al. (2020) write a slightly richer model: retrieval is a distribution, and you **marginalize** over documents,

$$p(y \mid q) = \sum_{z \in \mathcal{Z}} p_\eta(z \mid q) p_\theta(y \mid q, z).$$

In practice the sum is restricted to the top- k , and $p_\eta(z \mid q)$ is often a softmax over retrieval scores. You do not need to train p_η this week. You do need to know: **the original paper treats which document you fetched as a latent variable**; the lab treats it as a hard top- k list.

Recall@ k on a labeled item (question, gold chunk ids $\mathcal{G}$):

$$\text{recall @}k = \frac{|\{\text{retrieved ids}\} \cap \mathcal{G}|}{|\mathcal{G}|}.$$

If gold is `[d1]` and top-2 are `[d3], [d2]`, recall@2 is 0. Retrieval failed before the generator ran.

6. Positive points and negative points

Positive.

- Update a fact by editing a document and re-indexing. No full retrain.
- You can ask for citations and check them.
- Works with a frozen generator (LoRA is optional; that is RAFT / Week 10.3).
- Domain corpora (this course, a hospital, a legal archive) stay outside the weights.

Negative.

- If the right chunk is not in the top- k , generation cannot recover it. Garbage in, fluent garbage out.
- Extra latency: embed + search + generate.
- Chunking is a method. Too small: you lose the date. Too large: you retrieve noise. Overlap is a knob.
- **Lost in the middle:** gold sitting in a long stuffed list is often ignored. Raising k is not free.
- **Unfaithful citations:** the model uses [d1] and writes [d2]. Same bug class as unfaithful CoT (note 13.3).
- The index can be stale, biased, or confidential. Retrieval copies whatever you stored.
- Dense retrievers fail on rare tokens (ids, course numbers) unless you add keywords.

When not to use RAG. A closed-book puzzle that is only in the weights (arithmetic without a calculator; ReAct's tool is the better extra). A task that needs *behavior* change (stop being sycophantic): that is SFT / DPO, not an index.

7. What can fail

Four failure modes to name on the board.

1. **Miss.** Gold never entered the top- k . Measure recall@ k . Fix chunking, the embedder, k , or hybrid search.
2. **Near-miss.** A related chunk ranks high ([d3] "hours" for a closing-time question). The model over-trusts it.
3. **Ignore.** Gold is in the prompt; the model recites a pretrain snapshot anyway. That is a **reader** bug. Week 10.3 RAFT trains the reader. This week the reader is frozen.
4. **Unfaithful citation.** The used sentence is in [d1]; the model writes [d2].

ReAct (note 14.2) can treat retrieve as one tool among others, including a calculator.

A **no-retrieval baseline** is the parametric model on the same questions. If RAG does not beat it, the index or the prompt is the story, not "we added RAG."

8. Measurement

Report **retrieval recall@ k** on a labeled set of (question, gold chunk ids), then **answer accuracy** with an **attribution check**: does the cited id actually contain the claim? Always include a no-retrieval baseline. "We used RAG" is not a result.

9. Teaching this note

About **50 minutes** at the board, then **~10 minutes** of video. Students have not seen RAG as a method before. Do not start from cosine.

- **0–10 min.** What / why: frozen weights vs an editable library; RAG vs fine-tune vs dumping the corpus.
- **10–22 min.** Architecture: offline index (vector + text + id) and the online path. Write the cosine and top- k formulas.
- **22–34 min.** Worked three-doc cosine. $k = 1$ vs $k = 2$. Lost in the middle.
- **34–48 min.** Pros / cons. Four failures. Recall@ k and attribution. RAFT is training; this is inference.
- **Then** play Karpathy intro **27:43–33:32** (tool use: browser as retrieve-then-read) and mention **40:45–42:15** (custom GPTs / files as a private index). Pause when the model emits a search and reads the hits.

Lecture ideas follow **CS224N 2026 L10**. Original notes; that course is not copied.

10. Worked example


Three chunks, vocab {library, close, shuttle}, bag-of-words counts:

id	text (toy)	vector
d1	library close	(1, 1, 0)
d2	shuttle loop	(0, 0, 1)
d3	library hours	(1, 0, 0)

Query q = “library close”, vector (1, 1, 0). Cosine $\cos(a, b) = a^T b / (\|a\| \|b\|)$.

- $\cos(q, d1) = 2 / (\sqrt{2}\sqrt{2}) = 1$
- $\cos(q, d2) = 0$
- $\cos(q, d3) = 1 / (\sqrt{2} \cdot 1) = 1 / \sqrt{2} \approx 0.707$

Rank: d1, d3, d2. For $k = 1$ you stuff only d1 (right chunk for a closing-time question). For $k = 2$ you also stuff d3: recall can rise, but the generator now sees a near-miss about “hours” that may not include 23:00.



Lost in the middle: if the gold chunk sits in the middle of a long stuffed list, models often ignore it. That is a retrieval-plus-position bug, not “the model cannot read.” Keep k small in Lab 14.

If the model then cites [d2] while using the 23:00 fact from d1, that is unfaithful attribution.

11. Where students get stuck

- Treating RAG as a new neural net. It is a **conditioning** stack.
 - Measuring only answer accuracy. You need recall@k **and** a citation check.
 - Stuffing $k = 20$ and wondering why the model ignores the gold chunk.
 - Putting only vectors in the index and dropping the raw text / id, so you cannot cite.
 - Calling Week 10.3 RAFT “RAG.” RAFT **trains** the reader on retrieved bundles. This week the reader is frozen.
-

12. Video

Watch [Andrej Karpathy, Intro to Large Language Models](#), second half: **27:43–33:32** (browser / retrieval) and **40:45–42:15** (customization with files).

Pause on emit-search → read hits → write with sources. Same URL as notes **13.1**, **14.2**, and **15.1**. This note is the retrieve-then-generate picture; **14.2** is the tool loop. For the original paper’s sum-over-documents picture, skim Lewis et al., [arXiv:2005.11401](#), the RAG-sequence setup.

13. Practice

1. In one sentence each: what RAG is, and why a frozen LLM by itself is not enough for this semester’s late-policy PDF.

2. What belongs in the index besides the embedding vector? Why?
3. Write the cosine formula, then compute $\cos(q, d3)$ exactly as a fraction from the three vectors above. Write the $k = 2$ id list.
4. Why can raising k hurt the answer even if it raises recall@k?
5. You retrieve the right chunk and the answer is still wrong. Name two different bugs.
6. Gold chunk is d1. Retrieved top-2 are d3, d2. Recall@2 for this item? Answer cites [d1]. Attribution pass or fail?
7. Write Lewis's sum $p(y \mid q) = \sum_z p_\eta(z \mid q) p_\theta(y \mid q, z)$ in words. What does Lab 14 do instead of the sum?