

9.2 RLHF

Three stages. Only PPO samples the current policy. $\beta=0$ is reward hacking.

These notes match the lecture slides. Use **(slides)** on the course hub for the deck.

RLHF (InstructGPT / ChatGPT-style) is a three-stage stack: SFT, a reward model, then a **policy** optimized with RL (usually **PPO**) so sampled replies score high under the RM without wandering too far from SFT.

1. What RLHF is

RLHF is three stages, not one loss.

1. **SFT** on instruction pairs (note **8.1**). Call this π_{SFT} .
2. **RM** r_ϕ on chosen/rejected pairs (note **9.1**).
3. **PPO**: sample $y \sim \pi_\theta(\cdot \mid x)$, score with r_ϕ , update π_θ .

After SFT and an RM, you sample from the **current** policy, score with r , and update to raise reward minus KL to a frozen reference. Only stage 3 needs samples from the current policy. Stages 1–2 use a fixed dataset. If your code loads a CSV of completions and never calls `generate` inside the PPO loop, it is not on-policy RLHF.

PPO is on-policy: you need fresh rollouts from the current π_θ . That is expensive (generate, reward, critic). The **critic** (value head) estimates advantages so the policy gradient has lower variance. You do not need to derive PPO in this course; you need the picture and the KL term. InstructGPT is this stack at company scale; a seminar project usually stops at best-of- n or DPO.

2. Why we use it

SFT does not explore. It copies answers that were already in the file. Best-of- n uses r only at decode: the policy itself never learns. RLHF trains π to look high-reward **as it samples**, so the next draw is already better, not merely reranked.

PPO still exists because on-policy sampling can find replies that were never in the pair file. That is the conceptual parent of the week: preferences $\rightarrow$ a scalar of “better” $\rightarrow$ a policy that maximizes it with a leash. It works at company scale and is finicky in a seminar: sampling, KL, reward scale, and a moving policy. **DPO** (next note) skips the RM and the loop. Keep the three-box picture even if you never run PPO.


On-policy also means last week’s samples are stale: π_θ moved, so you generate again. That is the cost people are avoiding with DPO. For a project, running full PPO on an LLM is usually out of scope. Best-of- n with a tiny RM, or DPO on a small classifier, is enough to show you understood the stack.

3. Architecture

Three frozen-or-moving pieces sit together at the PPO stage:

- Policy π_θ (moving).
- Reference π_{ref} (usually SFT, **frozen**).
- RM r_ϕ (**frozen** in the PPO stage).
- Optional **critic** / value head for advantages.

RLHF: climb the reward, stay near the SFT policy (KL penalty)



Why $\pi_{\text{ref}} = \pi_{\text{SFT}}$ not the raw pretrained LM: the RM was trained on SFT-style answers. Leashing to the base model pulls you back to web-text, not to the assistant you already paid for. The three boxes are SFT, RM, and RL. **Best-of- n** is already a strong baseline in AlpacaFarm-style studies; it is not a fourth box in this loop.

4. How it works, step by step

1. Run SFT. Freeze that checkpoint as π_{SFT} and as π_{ref} .
2. Train the RM on pairs. Freeze r_ϕ .
3. **PPO loop.** Sample a batch of new y from the current π_θ . Score with r_ϕ . Estimate advantages with the critic. Update the policy with a clipped objective. Subtract a KL penalty toward π_{ref} .
4. **Read β .** Large β stays near SFT; small β chases reward. In code this is often a per-token penalty $\beta \log \pi_\theta - \beta \log \pi_{\text{ref}}$.
5. **If $\beta = 0$, there is no leash.** Expect reward to rise while KL explodes and the text becomes a caricature of the raters: long, flattering, or degenerate. That is reward hacking, not “learning faster.”

Without a constraint, π_θ **hacks** r_ϕ . The KL term is the method, not a footnote. You do not derive the PPO clip in this course; you need to know that a clip plus an advantage estimate is how the policy step is usually written, and that $\beta = 0$ is the failure mode to name.

5. Mathematical formulas

The usual RLHF objective looks like

$$\max_{\pi} \mathbb{E}[r(x, y)] - \beta \text{KL}(\pi(\cdot | x) \| \pi_{\text{ref}}(\cdot | x)),$$

or, sampling from the current policy,

$$\mathbb{E}_{x, y \sim \pi_{\theta}} [r_{\phi}(x, y) - \beta \text{KL}(\pi_{\theta}(\cdot | x) \| \pi_{\text{ref}}(\cdot | x))].$$

KL cartoon, two-word vocab $\{a, b\}$. $\pi_{\text{ref}} = (0.7, 0.3)$, $\pi_{\theta} = (0.99, 0.01)$:

$$\text{KL}(\pi_{\theta} \| \pi_{\text{ref}}) = 0.99 \log \frac{0.99}{0.7} + 0.01 \log \frac{0.01}{0.3} \approx 0.99 \cdot 0.347 + 0.01 \cdot (-3.40) \approx 0.31.$$

Penalty $\beta \cdot 0.31$. If $\beta = 0.1$, you subtract 0.031 from reward; if $\beta = 0$, you subtract nothing and π_{θ} can spike. If $\pi_{\theta} = \pi_{\text{ref}}$, KL is 0 and the leash is silent at that instant.

 A two-word spike: mass on $\langle \text{math xmlns=$

a goes 0.7 $\rightarrow$ 0.99" src="graphics/9.2-rlhf/kl-numeric.png" />

6. Positive points and negative points

Positive.

- Can beat SFT on preference evals because the policy learns while it samples, not only at decode.
- On-policy rollouts can find replies that were never in the pair file.
- The KL leash is an explicit knob between “stay an assistant” and “chase the RM.”
- Best-of- n remains a fair baseline on the same RM; PPO has to beat it.

Negative.

- Unstable and expensive: generate, reward, critic, and stale samples every time π_{θ} moves.
- RM misspecification becomes policy misspecification. Reward hacking is what $\beta = 0$ looks like.
- Using the pretrained base as π_{ref} after SFT already moved the policy pulls you back to web-text.
- Drawing RLHF as one loss hides that only stage 3 is RL.
- Full PPO on an LLM is usually out of scope for a seminar project, which is why people look to DPO.

When not to. If you have pairs and a GPU but no appetite for a PPO stack, skip to DPO (next note) or stop at best-of- n . If $\beta = 0$ “to learn faster,” you are not running RLHF; you are maximizing a misspecified RM.

7. Teaching this note

About **35 minutes** at the board: three boxes (SFT, RM, PPO), then write the KL objective and set $\beta = 0$ vs large β . Play the same Karpathy **RLHF** window **21:05–25:43** (and a minute after if he is still on comparisons). DPO is next; do not derive PPO.

8. Worked example

Stage check: a batch of instruction pairs $\rightarrow$ SFT. A batch of $(y_w, y_l) \rightarrow$ RM. A batch of **new** samples from $\pi_\theta \rightarrow$ PPO. If your code loads a CSV of completions and never calls `generate` inside the PPO loop, it is not on-policy RLHF.

KL cartoon: two distributions on a 2-word vocab $\{a, b\}$. $\pi_{\text{ref}} = (0.7, 0.3)$, $\pi_\theta = (0.99, 0.01)$.

$$\text{KL}(\pi_\theta \parallel \pi_{\text{ref}}) = 0.99 \log \frac{0.99}{0.7} + 0.01 \log \frac{0.01}{0.3} \approx 0.99 \cdot 0.347 + 0.01 \cdot (-3.40) \approx 0.31.$$

Penalty $\beta \cdot 0.31$. If $\beta = 0.1$, you subtract 0.031 from reward; if $\beta = 0$, you subtract nothing and π_θ can spike.

The three boxes are SFT, RM, and RL. **Best-of- n** is already a strong baseline in AlpacaFarm-style studies. PPO still exists because on-policy sampling can find replies that were never in the pair file. InstructGPT is this stack at company scale; a seminar project usually stops at best-of- n or DPO.

Why $\pi_{\text{ref}} = \pi_{\text{SFT}}$ not the raw pretrained LM: the RM was trained on SFT-style answers. Leashing to the base model pulls you back to web-text, not to the assistant you already paid for.

9. Where students get stuck

- Drawing RLHF as one loss. It is three stages; only the last is RL.
 - Setting $\beta = 0$ "to learn faster" and then being surprised by reward hacking.
 - Using the pretrained base as π_{ref} after SFT already moved the policy.
-

10. Video

[Karpathy — Intro to Large Language Models](#). Play **21:05–25:43** (RLHF, comparison labels). Pause on the three-stage story: write, compare, then a loop that uses those comparisons. This clip will not derive PPO; the board does the KL leash.

11. Practice

1. Which of the three stages needs *samples from the current policy*, not a fixed dataset?

2. If $\beta = 0$, what failure mode should you expect?
3. Why is π_{ref} the SFT model rather than the raw pretrained LM?
4. $\pi_{\theta} = (0.8, 0.2)$, $\pi_{\text{ref}} = (0.8, 0.2)$. What is $\text{KL}(\pi_{\theta} \parallel \pi_{\text{ref}})$, and what does that say about the leash at this instant?
5. Reward $r = 5$, $\beta = 2$, $\text{KL} = 1.5$. What scalar is the objective $r - \beta \text{KL}$? If you raise β to 4, who wins: the RM or the reference?