

7.1 Scaling Laws

Kaplan power laws, then Chinchilla: GPT-3 was 1.7 tokens/parameter, not 20.

These notes match the lecture slides. Use **(slides)** on the course hub for the deck.

Bigger models, trained on more data with more compute, tend to get better at next-token prediction.

Scaling laws make that trend a curve you can fit, not a slogan. They do not promise a better chatbot; they promise a lower **pretraining loss** if the recipe stays the same.

1. What a scaling law is

A **scaling law** is a fitted curve of pretraining **cross-entropy** versus scale. You plot loss against compute C , parameter count N , or dataset size D , and you ask whether the cloud of runs looks like a smooth power law plus a floor. It is an empirical regularity for a fixed architecture family and a fixed data mixture, not a theorem that the next chatbot will be more helpful.

Kaplan et al. (2020) fit language-model loss against compute (and against N and D) and saw smooth **power laws** over many orders of magnitude:

$$L(C) \approx aC^{-b} + L_{\infty}.$$

The irreducible term L_{∞} is entropy of the data plus whatever the architecture cannot represent. The interesting bit is the slope $-b$: extra FLOPs still buy loss, with diminishing returns. Once measured loss sits on L_{∞} , the curve has already told you that more compute will not move the number.

Hoffmann et al. (2022), **Chinchilla**, asked a different question with the same ingredients. At a **fixed FLOP budget**, how should you split compute between a larger net and a longer token stream? Their fit said you should scale **tokens and parameters together**, not dump the budget into a huge model that sees the data once. The rule of thumb from that paper is about **20 tokens per parameter** for the models they trained. Put it on the board as a rule of thumb from one paper, not a law of nature.

A scaling law is therefore two claims at once. Kaplan: loss versus scale looks like a power law plus a floor. Chinchilla: at fixed FLOPs, grow N and D together. Neither claim is “we trained a 7B.”

2. Why we use it

“We trained a 7B” is not a complete sentence. You need the token budget and whether you were compute-optimal. A smaller model trained longer can beat a larger undertrained one. That is why later LLaMA-style recipes look “small” next to GPT-3 and still read well: they were fed more tokens.

Early GPT-3-style runs were **compute-optimal for a short training budget** that underused data: a large N , not enough tokens. GPT-3 was 175 B parameters on 300 B tokens, about 1.7 tokens per



parameter. Chinchilla's thumb is nearer 20. A 70 B model trained with enough tokens beat much larger undertrained nets.


You use the curve to **plan a run**. Fix two of $\{N, D, C\}$ and the third is constrained. Kaplan-style runs often **oversize** N for a short C : the compute is spent on a wide net that barely sees D . Chinchilla responds that C on more tokens. Same FLOP budget, different (N, D) . Write both numbers, not only "7B."

A 7B model at 20 tokens/parameter wants about $7 \times 10^9 \times 20 = 1.4 \times 10^{11}$ tokens. If you only have 14B tokens, you are $10 \times$ short of that rule of thumb—not "we scaled." Lab 7 will fit $L = aC^{-b} + c$ on constructed points so you see the algebra, not a GPU cluster.

3. Architecture

This is not a new network. The object on the board is a **log-log plot** and a recipe (N, D, C) . The transformer is whatever you already had; the method is how you choose its size and its token budget.





On the Kaplan plot, the x-axis is compute and the y-axis is pretraining loss. The curve falls, then flattens toward L_∞ . On the Chinchilla picture, two models sit at the same FLOP budget: one is wide and starved of tokens; the other is smaller and fed more data. The second is the balanced pair. Training FLOPs scale roughly as $6ND$ (forward plus backward), so holding C fixed really does mean holding the product ND fixed while you trade N against D .

4. How it works, step by step

You do not train a new layer. You fit a curve, then you do arithmetic that makes “7B” a pair of numbers.

1. **Collect runs** (or, in Lab 7, constructed points) of pretraining loss versus compute. The range in the papers is many orders of magnitude; four noisy toy points are for the algebra.
2. **Fit** $L \approx aC^{-b} + L_\infty$. Read b as diminishing returns and L_∞ as the floor you cannot train through.
3. **Name the budget.** Training FLOPs cartoon: $\approx 6ND$. If you double N and halve D , C is unchanged.
4. **Apply Chinchilla arithmetic.** A 7B model at 20 tokens/parameter wants $7B \times 20 = 140B$ tokens. GPT-3’s 175B on 300B tokens is about 1.7 tokens per parameter: oversized N for that C .
5. **Write both numbers** in a report: checkpoint size and token budget. Two extra epochs on a toy set is not a scaling law.

Same FLOPs, two recipes. Model A: $N = 4N_0$, $D = D_0 / 4$. Model B: $N = N_0$, $D = D_0$. Same ND . Chinchilla picks the **balanced** pair, not the $4 \times$ -larger starved model. If the rule of thumb is 20 tokens/parameter, B matches it when $D_0 = 20N_0$; A has only 5 tokens/parameter.

5. Mathematical formulas

Kaplan's compute law:

$$L(C) \approx aC^{-b} + L_{\infty}.$$

Lab 7 writes the same fit as $L = aC^{-b} + c$ on constructed points. The slope $-b$ is diminishing returns; L_{∞} (or c) is the floor.

Training FLOPs, cartoon:

$$C \approx 6ND.$$

Chinchilla's planning thumb, from that paper's fit, is about 20 tokens per parameter:

$$D \approx 20N.$$

A 7B model then wants $7 \times 10^9 \times 20 = 1.4 \times 10^{11}$ tokens. GPT-3's ratio was $300B / 175B \approx 1.7$. You do not need the papers' full tables of a and b this week. You need to write $L(C)$, mark L_{∞} , and compute a token budget at fixed N .

6. Positive points and negative points

Positive.

- A scaling law turns "bigger is better" into a curve you can fit and a recipe (N, D, C) you can plan against.
- Chinchilla arithmetic stops "we trained a 7B" from being a complete sentence: you also write the token budget.
- Same FLOP budget, different (N, D) : a smaller model trained longer can beat a larger undertrained one.
- Lab 7's toy fit shows the algebra of $aC^{-b} + c$ without a cluster.

Negative.

- L is pretraining cross-entropy, not chat quality. Downstream accuracy is not L .
- Alignment (Weeks 8–9), data quality, mixture, and context length all sit off the curve.
- Inference cost still scales with N even if you already paid for training.
- Four noisy toy points are not a law. The papers' range is many orders of magnitude.
- The 20 tokens/parameter figure is a rule of thumb from one paper, not a law of nature.

When not to. Do not claim "we scaled" because you trained two extra epochs on a toy set. If your measured loss already sits on L_{∞} , extra FLOPs buy almost nothing; change data or architecture instead of buying more compute.

7. Teaching this note

About **35 minutes** at the board: write $L(C)$, mark L_∞ , then a Chinchilla arithmetic example (tokens vs parameters at fixed FLOPs). Play Karpathy **after 20:00**, especially **25:43–27:43** (scaling laws, ~2 min, plus a little context). The fit in Lab 7 is later in the week.

Put 20 tokens/parameter on the board as a **rule of thumb from one paper**, not a law of nature. Then do the $7B \times 20$ arithmetic so “we trained a 7B” is never a complete sentence.

8. Worked example

Suppose $L = 2 C^{-0.15} + 1.4$. At $C = 10^4$,

$$L \approx 2 \cdot 10^{-0.6} + 1.4 \approx 2 \cdot 0.251 + 1.4 \approx 1.90.$$

At $C = 10^6$, $10^{-0.9} \approx 0.126$, so $L \approx 2 \cdot 0.126 + 1.4 = 1.65$. A $100\times$ compute jump bought 0.25 nats, not a new architecture. If your measured loss is already 1.42, you are sitting on L_∞ : extra FLOPs buy almost nothing.

Chinchilla, same FLOPs. Training FLOPs scale roughly as $6ND$ (forward+backward). Model A: $N = 4N_0$, $D = D_0 / 4$. Model B: $N = N_0$, $D = D_0$. Same ND . Chinchilla picks the **balanced** pair, not the $4\times$ -larger starved model. If the rule of thumb is 20 tokens/parameter, B matches it when $D_0 = 20N_0$; A has only 5 tokens/parameter.

9. Where students get stuck

- Reading L as “chat quality.” It is pretraining cross-entropy.
 - Thinking a bigger N always wins at fixed FLOPs. Chinchilla is the counterexample.
 - Fitting a, b, c on four noisy toy points and claiming a law. Lab 7 is the algebra; the paper’s range is many orders of magnitude.
-

10. Video

[Karpathy — Intro to Large Language Models](#). Play from **20:00**, and pause on **25:43–27:43** (scaling laws). The first 20 min were Week 1; this hour is the compute curve and why tokens matter as much as parameters.

11. Practice

1. Two models, same FLOPs. One has $4\times$ the parameters and $1/4$ the tokens of the other. Who does Chinchilla pick, and why?

2. If L_∞ is already almost your measured loss, what does extra compute buy?
3. Why can a scaling law look great while a chat eval is flat?
4. Using $L = 2C^{-0.15} + 1.4$, compute L at $C = 10^5$. ($10^{-0.75} \approx 0.178$.)
5. A 1B-parameter model trained on 4B tokens. Using 20 tokens/parameter, is it over- or under-trained on data? By what factor?