

# Course Map and the Semester Project

DATA 443/643 · Advanced Concepts in Large Language Models

Week 1

American University

Ahmad Mousavi

Ideas and 16:9 layout from Stanford CS224N, CS231N, and MIT 6.S191. Original slides; those courses are not copied.

# Lecture plan


- 1 Where this course sits, and why it is self-contained
- 2 An LLM is a stack; Week 1 is the base
- 3 What pretraining actually minimizes
- 4 The 140 GB weights file
- 5 Turn a method into a project question

Key goal: leave with a stack picture, a next-token loss, and one measurable project sentence.

# A self-contained LLM course

- Listed prereq: DATA 427/627.
- DATA 641 and 642 are not required.
- Overlap is intentional so more people can finish a project.


Where Week 1 sits: foundations before transformers



# A language model is a stack

- Chat products sit at the top.
- The weights still came from next-token training.
- Name the layer you will change: data, prompt, adapter, retriever, or eval.

A language model is a stack. Week 1 is the bottom two layers.



# Pretraining minimizes next-token NLL

$$P(w_1, \dots, w_T) = \prod_{t=1}^T P(w_t \mid w_{1:t-1})$$

Each factor is a softmax over the vocabulary.

The training loss is the negative log of those probabilities.

If the model assigns 0.70, NLL is about 0.36. If 0.10, NLL is about 2.30.

# Llama 2 70B is a 140 GB file

$$70 \times 10^9 \times 2 \text{ bytes} = 140 \text{ GB}$$

Each parameter stored as float16 (2 bytes).

Parameters file vs run file: running needs extra memory.

Karpathy: an LLM is weights plus a little code.

# The project starts with a question

## Not a project

Use LoRA on news.

## A project

Does a LoRA adapter on local news reduce entity hallucination vs. the base model, measured by exact-match on a 100-item holdout?

# Today's two-hour meeting

## 1.1 Course map

Stack, next-token NLL, 140 GB, project sentence. Short Karpathy clip.

## 1.2–1.4 Board notes

Neuron and XOR, then gradient descent, then skip-gram and bias.

## Lab 1

Linear vs MLP on XOR, then the constructed 2-D embedding probe.

Key goal: one timed meeting, one stack picture, one project sentence.