


The bottleneck

- An RNN compresses the past into one vector.
- Attention: every token looks at every token in one step.
- Cost: $O(T^2)$. Path length: 1.

RNN: each state sees the past through a chain



Attention: every token can look at every token

