


Attention has no order

- Add sinusoidal or learned positions.
- Later models use relative or rotary positions.



One block

- Attention mixes across positions. The MLP mixes across channels.
- Residuals are the cousin of the LSTM highway.

One transformer block

