


Patches as tokens

- Flatten each patch, map to width d , add positions.
- A 224 image with patch 16 is 196 tokens: a short paragraph.

A ViT treats patches the way an LLM treats tokens



The same stack as language

Vision transformer: same block as language, different tokens

- [CLS] + positions + transformer.
- In this course, start from a pretrained ViT, not from scratch.

