

Two towers, one cosine

- Image encoder and text encoder, web-scale pairs.
- Good at retrieval and zero-shot. Not a captioner.

CLIP: two towers, one cosine, a batch of matches

