

Scaling Laws

DATA 443/643 · Week 7, note 7.1

Week 7

American University

Ahmad Mousavi

Ideas from Stanford CS224N 2025 L9 (Kaplan curves; GPT-3 tokens vs Chinchilla). Original slides; that course is not copied.


Lecture plan

- 1 Loss as a power of compute
- 2 A numeric on $L(C)$
- 3 Chinchilla: tokens with parameters
- 4 What a law does not tell you

Key goal: write N and D , not only $7B$. L is pretraining CE, not chat quality.

A smoother log-log curve

- $L(C) \approx aC^{-b} + L_{\infty}$.
- Extra FLOPs still buy loss, with diminishing returns.
- Lab 7 fits a , b , c on toy points. The paper's range is many decades.



One hundred times the compute

$$L = 2 C^{-0.15} + 1.4$$

$$C = 10^4: L \approx 1.90.$$


$$C = 10^6: L \approx 1.65. \text{ A}$$

100× jump bought 0.25 nats.

If you already sit on L_∞ , extra FLOPs buy almost nothing.


GPT-3 was underfed

- CS224N L9: 175B parameters, 300B tokens.
- That is about 1.7 tokens per parameter, not 20.
- A 70B trained with enough tokens beat larger starved nets.



Same FLOPs, different (N, D)

- Training FLOPs scale like $6ND$.
- A: $4N_0$ params, $D_0 / 4$ tokens. B: N_0 , D_0 .
- Chinchilla picks the balanced pair, not the wide starved net.



What to take from this lecture

Write D

7B × 20 tokens/param wants 140B tokens.

L is not chat

Alignment, mixture, and eval sit off the curve.

Lab 7

Fit a, b, c. Four noisy points are not a law.