


Loss versus compute

- On a log–log plot, pretraining loss falls as a power of compute.
- The curve bends when you starve data or parameters.



Chinchilla: feed the model tokens

- Bigger is not enough if tokens do not keep up.
- A compute-optimal run balances parameters and data.

