

Mixture of Experts

DATA 443/643 · Week 7, note 7.2

Week 7

American University

Ahmad Mousavi

Ideas from Stanford CS336 2025 L4 (Mixtral: 8 routed, top-2 active). Original slides; that course is not copied.

Lecture plan


- 1 Router, then k experts
- 2 Stored weights versus active FLOPs
- 3 Load collapse
- 4 When a dense 7B is the better project

Key goal: 8 MLPs on disk, 2 in the FLOP count. Say that out loud.

Sparse where the parameters already live

- Attention can stay dense. The MLP is the MoE.
- CS336: Mixtral is 8 experts, $k=2$. Switch used $k=1$.
- DeepSeek's 256-expert nets wait. This hour is the picture.

Sparse MoE: most experts stay off for a given token



Top-k mix

$$y = \sum_{i \in \mathcal{T}(x)} g_i(x) E_i(x)$$


$\mathcal{T}$ is the k selected experts.

Renormalize the two gates so they sum to 1.

If every token has the same logits, you trained two MLPs.

Collapse is a dense net in disguise

- 32 tokens, $k=2$: 64 slots. Uniform load is 8 per expert.
- If two experts eat 43 slots, six experts are cold.
- Load-balancing tries to stop that. Token drop is the other knob.



What you actually ship

Dense 7B

One weight file, a dense GEMM, simple batching.

Mixtral-style 8x

Eight MLP shards in RAM even if two run. Tokens in one sequence may disagree on the expert.

What to take from this lecture

Stored $\neq$ active

8× MLP weights, $\sim 2\times$ MLP FLOPs.

Router

If it always picks expert 2, you trained one expert.

Lab 7

Scaling fit and INT8, not an MoE train.