

## A router, then a few experts

- Each token activates a sparse subset of MLPs.
- Capacity grows without dense FLOPs on every expert.

Sparse MoE: most experts stay off for a given token

