


# Prune, quantize, distill

- Compression is a pipeline. Measure quality after each cut.
- Int8 is not free: report quality, not only size.

Compression is a pipeline. Measure quality after each cut.



# Speculative decoding

- A small model drafts a prefix. The large model verifies.
- Accepted tokens must match the large model's distribution.

Speculative decoding: guess many tokens, accept the prefix the big model agrees on

