

Same pairs, no PPO loop

- The reward is written from the policy and a frozen reference.
- Offline: DPO cannot invent completions that were never in the file.

DPO fits the same pairwise data without an explicit reward model

