

Safety as a Rate

DATA 443/643 · Week 10, note 10.1

Week 10

American University

Ahmad Mousavi

Ideas from Stanford CS336 2025 L12 (HarmBench / HELM safety as rates). Original slides; that course is not copied.

Lecture plan


- 1 A fixed probe list, then a fraction
- 2 Split harm from over-refusal
- 3 Rerun after the patch
- 4 A judge is not a substitute for a sample

Key goal: 4/40 is a number you can rerun. A vibe is not.

Probes in, a log out


- CS336 L12: HarmBench is 510 labeled behaviors.
- AIR-Bench: 314 categories, 5694 prompts. HELM hosts suites.
- In class you use a short public list and instructor categories.

Red-teaming is an evaluation, not a vibe. Log attacks that work.



Four cells, one rate

- 10 + 10 + 10 + 10 probes. Fails 0, 1, 2, 1.
- Overall $4 / 40 = 10\%$. Log the four ids.
- Do not publish working attack strings on the open site.



Two different failures

True safety hit

A disallowed or leaky completion. Count it in the harm cell.

Over-refusal

"What is ibuprofen used for?" should not refuse. Mixing this into "unsafe %" hides a model that tightened too far.

CS336: capability versus propensity

API model

Refusal rate (propensity) is the number you ship. The user cannot fine-tune the weights.

Open-weight

A refusal is not a capability bound. Fine-tuning can move the propensity. Report both.

What to take from this lecture

Rerun

Same 40 after DPO. One cell moved is not “done.”

Judge

Spot-check a 10% sample. CS224N L12: look at the outputs.

Not a recipe

Measurement. No jailbreak how-to in this course.