

2.2 Vanishing and Exploding Gradients

BPTT sums copies of the shared matrix. If the largest $|\text{eigenvalue}|$ is below 1, early tokens stop training.

These notes match the lecture slides. Use **(slides)** on the course hub for the deck.

Backprop through an unrolled RNN multiplies one Jacobian per time step. If those matrices shrink, the early-token gradient **vanishes**. If they grow, it **explodes**. Either way, long-range learning dies.

1. Backprop through time

The loss at each step depends on the shared W_h . The multivariable chain rule says: **the gradient with respect to a repeated weight is the sum of the gradient at each copy.**

 BPTT sums the copies of W_h

Wh" src="graphics/2.2-vanishing-gradients/bptt.png" />

That algorithm is **backpropagation through time** (Werbos 1988). In practice people often **truncate** after a few tens of steps so a batch fits in memory. Truncation is an efficiency choice. It does not fix vanishing: a truncated window never sees the token that sat 200 steps back.


You do not implement truncated BPTT by hand this week. `loss.backward()` on `nn.RNN` is BPTT over the sequence you fed it.

2. A product along the chain

The loss at the last token, $L(h_T)$, depends on h_1 through

$$\frac{\partial L}{\partial h_1} = \frac{\partial L}{\partial h_T} \frac{\partial h_T}{\partial h_{T-1}} \cdots \frac{\partial h_2}{\partial h_1}.$$

Each factor $\partial h_t / \partial h_{t-1}$ contains W_h and $\tanh'$. For $\tanh$, $|\tanh'| \leq 1$, and it is $\ll 1$ unless the pre-activation sits near 0. CS231N's slogan: multiplying by the activation Jacobian is **almost always a shrink**. A product of twenty numbers smaller than 1 is tiny.




The gated curve is the slogan for LSTM/GRU: keep a path whose multiplier can stay near 1.

In 1-D, if each factor is 0.5, ten steps give $0.5^{10} \approx 0.001$. The early token still sits in the forward state as a faint trace, but the **learning signal** is gone.

3. Eigenvalues, not just $\tanh$

Ignore the nonlinearity for a minute. A linear recurrence $h_t = W_h h_{t-1}$ scales, for large t , like $|\lambda_{\max}|^t$, where $\lambda_{\max}$ is the largest-magnitude eigenvalue of W_h (CMU 11-785 L14). CS231N says the same thing with the largest **singular value**.

- $|\lambda_{\max}| < 1$: the forward state and the backward signal **vanish**.
- $|\lambda_{\max}| > 1$: they **explode**.
- Exactly 1 in every direction is the unitary dream. Vanilla RNNs do not learn that matrix.




ReLU is not a rescue: if the pre-activation is negative the unit dies; if W_h is large it blows up. Tanh is the least-bad vanilla default. It is still short memory.

This is not only an RNN problem. Any deep stack of Jacobians can vanish or explode. Recurrence is just the case where you multiply by the **same** W_h over and over (Bengio et al. 1994; Pascanu et al. 2013).

4. Vanishing versus exploding


Same product, opposite problem.



If $\|W_h\|$ is large, the signal blows up, weights become NaNs, and training stops. **Gradient clipping** (Pascanu et al. 2013): if $\|g\| > c$, replace g by $c g / \|g\|$. Same direction, shorter step.

Exploding: same direction, shorter step

Vanishing: there is nothing to rescale



Clipping is the bandage for exploding. It does **not** fix vanishing. A vanished vector is already near 0; multiplying it by 1 does nothing useful.

5. What this means for language

A vanilla RNN can learn “the next word looks like the last few.” It struggles with “the negation was at the start of the review.”

CS224N’s tickets example, shortened: she tried to print her *tickets*, left, bought toner, came back, and finally printed her *_*. The model has to bind the last word to *tickets* many steps earlier. If the gradient from that last NLL never reaches the early embedding, the net never learns the dependency, and at test time it cannot use it.

In a project, if you still train an RNN on long documents, watch the gradient norm **and** (if you can) the contribution from early positions. If it is always ≈ 0 on the start of the sequence, the net is not using the beginning.

Lab 2 is a copy task: the first bit must survive many steps. Vanilla RNNs fail as the gap grows; that is this product, not a PyTorch bug.

Transformers side-step the chain: every token attends to every other token in one step (Week 3). Gates (Week 2.3) were the fix *inside* the recurrent family.

6. Teaching this note

~18 minutes. BPTT as a **sum over copies**, then the product of Jacobians, then the 1-D geometric example (0.5^T vs 1.5^T), then $|\lambda_{\max}|^t$, then clipping as a cap on a tall arrow. Play the vanishing/exploding segment of CS224N (about **45:00–58:00** on the linked older lecture; if chapter marks differ, jump to the Jacobian-product board). LSTM details wait for 2.3.

7. Worked example

Scalar RNN: $h_t = \tanh(wh_{t-1})$. Then $h'_t(h_{t-1}) = \tanh'(z_t)w$ with $|\tanh'| \leq 1$.

Take $w = 0.5$ and pretend $\tanh' = 1$ (best case). The factor per step is 0.5:

$$0.5^5 = 0.03125, \quad 0.5^{10} \approx 0.0010, \quad 0.5^{20} \approx 9.5 \times 10^{-7}.$$

Now $w = 2$: $2^{10} = 1024$. Same chain, opposite disaster.

Linear picture: $0.9^{20} \approx 0.12$, $1.1^{20} \approx 6.7$. Twenty steps already separate “almost gone” from “several times larger.”

Clipping: gradient $g = 50$, threshold $c = 5$. New gradient 5. If instead $g = 10^{-6}$, clipping leaves it at 10^{-6} .

Vector clip: $g = [-6, 8]$ has Euclidean norm 10. With $c = 1$, the clipped vector is $g / 10 = [-0.6, 0.8]$.

8. Where students get stuck

- Thinking clipping “fixes RNNs,” including vanishing.
 - Forgetting $\tanh' \leq 1$ and blaming only W_h .
 - Treating truncated BPTT as a vanishing-gradient solution.
 - Plotting loss and missing a **per-position** gradient plot in Lab 2.
-

9. Video

Watch [Stanford CS224N Lecture 6 \(RNNs / vanishing gradients\)](#).

Play the vanishing and exploding segment (roughly **45:00–58:00**). Pause on the product of Jacobians and on the plot of gradient size vs. steps back in time. Skip n-gram language modeling and skip NMT.

The matching *form* lectures this week are CS224N W26 L4 (near vs. long-term effects, clipping) and CS231N 2025 L7 (computational graph, singular values). CMU 11-785 L14 is the eigenvalue/memory hour that 2.3 continues.

10. Practice

1. Why is $\tanh'$ a problem when the state saturates at ± 1 ?
2. Clipping stops explosions. Why does it not restore a vanished signal?
3. In one sentence: what would you plot in TensorBoard to see this failure in Lab 2?
4. If each Jacobian factor has operator norm 0.8, the bound on $\|\partial h_{20} / \partial h_1\|$ is 0.8^{19} . Compute 0.8^{10} and 0.8^{19} (calculator is fine). Is the 19-step bound closer to 0.1 or to 0.01?
5. A clip threshold $c = 1$ meets a gradient of $[-6, 8]$ (Euclidean norm 10). Write the clipped vector.
6. A linear RNN has $|\lambda_{\max}(W_h)| = 0.9$. After 30 steps, is $|\lambda_{\max}|^{30}$ closer to 0.04 or to 0.4?

7. Why does summing gradients over time copies of W_h not, by itself, save an early token whose Jacobian product is 10^{-7} ?